



YOLO Assignment

MTech in Applied AI
Deep Learning Techniques

by

Sanjeev Kumar Pandey*
Faculty: Dr Anamika Gupta

Submitted May 21, 2025

* Student ID: 31050, Enrolment No:MT23AAI001, Email: pandey.sanjeev@yahoo.com

Table of Contents

1	Programming Assignment	2
1.1	Question 1	2
1.2	Question 2	3
1.3	Question 3	4
2	Theory Assignment	34
2.1	Question 1	34
2.2	Question 2	36
3	End of Document	38

1 Programming Assignment

1.1 Question 1

Define functions to convert between these two representations: corner_to_center converts from the two-corner representation to the center-width-height presentation, and center_to_corner vice versa.

1.1.1 Solution: Python code

Here are the two functions to convert bounding boxes between the two common YOLO representations:

1. `corner_to_center`: Converts from the two-corner format

$$(x_{\min}, y_{\min}, x_{\max}, y_{\max})$$

to the center-width-height format

$$(x_{\text{center}}, y_{\text{center}}, \text{width}, \text{height}).$$

2. `center_to_corner`: Converts from the center-width-height format back to the two-corner format.

```

1 def corner_to_center(x_min, y_min, x_max, y_max):
2     """
3         Convert bounding box from corner format (x_min, y_min, x_max, y_max)
4         to center format (x_center, y_center, width, height).
5     """
6     width = x_max - x_min
7     height = y_max - y_min
8     x_center = x_min + width / 2
9     y_center = y_min + height / 2
10    return x_center, y_center, width, height
11
12
13 def center_to_corner(x_center, y_center, width, height):
14     """
15         Convert bounding box from center format (x_center, y_center, width, height)
16         to corner format (x_min, y_min, x_max, y_max).
17     """
18     x_min = x_center - width / 2
19     y_min = y_center - height / 2
20     x_max = x_center + width / 2
21     y_max = y_center + height / 2
22     return x_min, y_min, x_max, y_max
23
24
25 corner_box = (100, 150, 200, 250)
26 center_box = corner_to_center(*corner_box)
27 corner_box_converted_back = center_to_corner(*center_box)
28 print(f'Corner to center : {center_box}')
29 print(f'Center to corner : {corner_box_converted_back}')

```

Listing 1: Code

1.1.2 Output Question 1

```
C:\Sanjeev\VNIT_CLASSES\VUNIT-AAI-SEM3\subjects\com\anamika_singh\YOL01_Assignment>python programming_assignment_q1.py
Corner to center : (150.0, 200.0, 100, 100)
Center to corner : (100.0, 150.0, 200.0, 250.0)
```

Figure 1: Output: Question 1

1.2 Question 2

Find an image and try to label a bounding box that contains the object.

1.2.1 Input Image



Figure 2: Input Image: Question 2

```
Microsoft Windows [Version 10.0.26100.3915]
(c) Microsoft Corporation. All rights reserved.

C:\Users\urssa>pip install opencv-python
Collecting opencv-python
  Downloading opencv_python-4.11.0.86-cp37abi3-win_amd64.whl.metadata (20 kB)
Requirement already satisfied: numpy>=1.21.2 in c:\users\urssa\appdata\local\programs\python\python311\lib\site-packages (from ope
ncv-python) (1.24.3)
  Downloading opencv_python-4.11.0.86-cp37abi3-win_amd64.whl (39.5 MB)
    39.5/39.5 MB 4.5 MB/s eta 0:00:00
Installing collected packages: opencv-python
Successfully installed opencv-python-4.11.0.86
C:\Users\urssa>
```

Figure 3: Depedency Installation

1.2.2 Solution: Python code

```
1 # Find an image and try to label a bounding box that contains the object.
2
3
4
5 import cv2
6
7 # Load the image
8 image = cv2.imread('black_apples.jpeg')
9
10 # Define the bounding box for the center apple (x_min, y_min, x_max, y_max)
11 # These values are estimated for the center apple in your image
12
13 y_min = 40
14 y_max = 310
15
16 x_min = 10
17 x_width = 320
18
19 for i in range(1, 4):
20     if i>1:
```

```

21     x_min = x_min + x_width + 20
22     x_max = x_min + x_width - 25
23
24     # Draw the bounding box
25     cv2.rectangle(image, (x_min, y_min), (x_max, y_max), color=(0, 255, 0), thickness=2)
26
27     # Add the label above the bounding box
28     label = "black apple"
29     cv2.putText(image, label, (x_min, y_min - 10), cv2.FONT_HERSHEY_SIMPLEX,
30                 fontScale=0.8, color=(0, 255, 0), thickness=2)
31
32 # Save or display the result
33 cv2.imwrite('labeled_black_apple.jpg', image)
34 # To display the image (optional):
35 # cv2.imshow('Labeled Image', image)
36 # cv2.waitKey(0)
37 # cv2.destroyAllWindows()

```

Listing 2: Code

1.2.3 Output Image



Figure 4: Input Image: Question 2

1.3 Question 3

Perform the following tasks in Python

1. Setup and Load Model

- (a) Install the required libraries.
- (b) Load a pre-trained YOLOv3 model.

2. Inference on Images

- (a) Run the model on sample images.
- (b) Visualize the bounding boxes with labels and confidence scores.

3. Custom Dataset Inference

- (a) Use a small custom dataset (e.g., traffic signs, animals, faces).
- (b) Run inference and evaluate the model's performance (precision, recall, mAP).

```
1 C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\subjects\com\anamika_singh\YOLO1_Assignment>pip
2 Collecting ultralytics
3   Downloading ultralytics-8.3.139-py3-none-any.whl.metadata (37 kB)
4 Requirement already satisfied: numpy>=1.23.0 in c:\users\urssa\appdata\local\programs\
5 Requirement already satisfied: matplotlib>=3.3.0 in c:\users\urssa\appdata\local\progr
6 Requirement already satisfied: opencv-python>=4.6.0 in c:\users\urssa\appdata\local\pr
7 Requirement already satisfied: pillow>=7.1.2 in c:\users\urssa\appdata\local\programs\
8 Requirement already satisfied: pyyaml>=5.3.1 in c:\users\urssa\appdata\local\programs\
9 Requirement already satisfied: requests>=2.23.0 in c:\users\urssa\appdata\local\program
10 Requirement already satisfied: scipy>=1.4.1 in c:\users\urssa\appdata\local\programs\p
11 Requirement already satisfied: torch>=1.8.0 in c:\users\urssa\appdata\local\programs\p
12 Collecting torchvision>=0.9.0 (from ultralytics)
13   Downloading torchvision-0.22.0-cp311-cp311-win_amd64.whl.metadata (6.3 kB)
14 Requirement already satisfied: tqdm>=4.64.0 in c:\users\urssa\appdata\local\programs\p
15 Collecting psutil (from ultralytics)
16   Downloading psutil-7.0.0-cp37-abi3-win_amd64.whl.metadata (23 kB)
17 Collecting py-cpuinfo (from ultralytics)
18   Downloading py_cpuinfo-9.0.0-py3-none-any.whl.metadata (794 bytes)
19 Requirement already satisfied: pandas>=1.1.4 in c:\users\urssa\appdata\local\programs\
20 Collecting ultralytics-thop>=2.0.0 (from ultralytics)
21   Downloading ultralytics_thop-2.0.14-py3-none-any.whl.metadata (9.4 kB)
22 Requirement already satisfied: contourpy>=1.0.1 in c:\users\urssa\appdata\local\program
23 Requirement already satisfied: cycler>=0.10 in c:\users\urssa\appdata\local\programs\p
24 Requirement already satisfied: fonttools>=4.22.0 in c:\users\urssa\appdata\local\progr
25 Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\urssa\appdata\local\progr
26 Requirement already satisfied: packaging>=20.0 in c:\users\urssa\appdata\local\program
27 Requirement already satisfied: pyparsing<3.1,>=2.3.1 in c:\users\urssa\appdata\local\p
28 Requirement already satisfied: python-dateutil>=2.7 in c:\users\urssa\appdata\local\pr
29 Requirement already satisfied: pytz>=2020.1 in c:\users\urssa\appdata\local\programs\p
30 Requirement already satisfied: tzdata>=2022.1 in c:\users\urssa\appdata\local\programs
31 Requirement already satisfied: six>=1.5 in c:\users\urssa\appdata\local\programs\pytho
32 Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\urssa\appdata\local
33 Requirement already satisfied: idna<4,>=2.5 in c:\users\urssa\appdata\local\programs\p
34 Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\urssa\appdata\local\prog
35 Requirement already satisfied: certifi>=2017.4.17 in c:\users\urssa\appdata\local\prog
36 Requirement already satisfied: filelock in c:\users\urssa\appdata\local\programs\python
37 Requirement already satisfied: typing-extensions>=4.10.0 in c:\users\urssa\appdata\loc
38 Requirement already satisfied: sympy>=1.13.3 in c:\users\urssa\appdata\local\programs\
39 Requirement already satisfied: networkx in c:\users\urssa\appdata\local\programs\python
40 Requirement already satisfied: jinja2 in c:\users\urssa\appdata\local\programs\python\
41 Requirement already satisfied: fsspec in c:\users\urssa\appdata\local\programs\python\
42 Requirement already satisfied: mpmath<1.4,>=1.1.0 in c:\users\urssa\appdata\local\pro
43 Requirement already satisfied: colorama in c:\users\urssa\appdata\local\programs\python
44 Requirement already satisfied: MarkupSafe>=2.0 in c:\users\urssa\appdata\local\program
45 Downloading ultralytics-8.3.139-py3-none-any.whl (1.0 MB)
46
47 Downloading torchvision-0.22.0-cp311-cp311-win_amd64.whl (1.7 MB)
48
49 Downloading ultralytics_thop-2.0.14-py3-none-any.whl (26 kB)
50 Downloading psutil-7.0.0-cp37-abi3-win_amd64.whl (244 kB)
51 Downloading py_cpuinfo-9.0.0-py3-none-any.whl (22 kB)
52 Installing collected packages: py-cpuinfo, psutil, ultralytics-thop, torchvision, ultr
```

```
53 Successfully installed psutil-7.0.0 py-cpuinfo-9.0.0 torchvision-0.22.0 ultralytics-8.0.0
```

Listing 3: Install the required libraries for YOLOv8 Model

1. Inference on Images

- (a) Run the model on sample images.
- (b) Visualize the bounding boxes with labels and confidence scores.

1.3.1 Input Image



Figure 5: Input Image: Question 3 a



Figure 6: Input Image: Question 3 b

```

1  from ultralytics import YOLO
2  import glob
3  import pandas as pd
4  # Load the pre-trained YOLOv8 model
5  model = YOLO('yolov8n.pt')
6
7  # List of image paths
8  image_paths = glob.glob('q3_2_*.jpeg')
9
10 # Run inference and save results to results/yolo_inference/
11 results = model.predict(
12     image_paths,
13     save=True,                      # Save images with bounding boxes
14     project='results',              # Top-level results folder
15     name='yolo_inference',          # Subfolder for this run
16     exist_ok=True                   # Overwrite if the folder exists
17 )
18
19 for result in results:
20     # Show the results in a window (requires GUI support)
21     result.show()
22
23
24 # To print detected boxes, labels, and confidence scores:
25 for result in results:
26     boxes = result.boxes
27     for box in boxes:
28         xyxy = box.xyxy.cpu().numpy()[0]    # bounding box coordinates (x_min, y_min, x_max, y_max)
29         conf = box.conf.cpu().numpy()[0]      # confidence score
30         cls = int(box.cls.cpu().numpy()[0])   # class index
31         label = model.names[cls]            # class label
32         print(f"Label: {label}, Confidence: {conf:.2f}, Box: {xyxy}")

```

Listing 4: Install the required libraries for YOLOv8 Model

1.3.2 Output Image



Figure 7: Output Image: Question 3 a

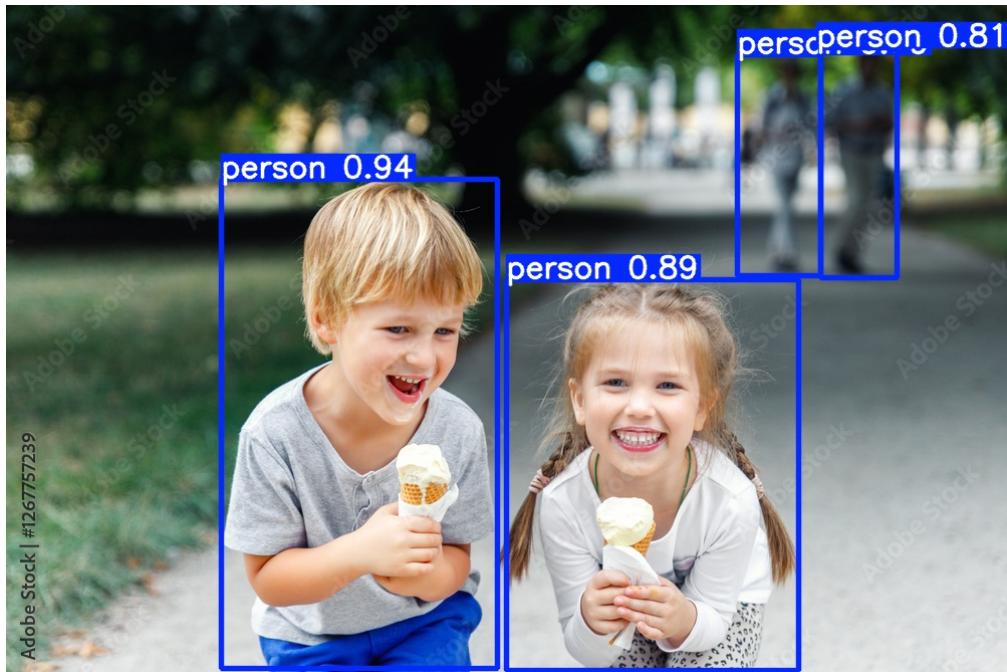


Figure 8: Output Image: Question 3 b

3. Custom Dataset Inference

1. Use a small custom dataset (e.g., traffic signs, animals, faces).
2. Run inference and evaluate the model's performance (precision, recall, mAP).

Automate YOLO Compatible Label

This script automates the creation of YOLO-compatible label files from images by leveraging a pretrained YOLOv8 model to detect objects and generate bounding box annotations. These label files are essential for training custom YOLO models, as they provide the ground truth bounding boxes and class labels in the required format.

Key Points in the Script

Bounding Box Conversion

The bounding box coordinates (xyxy) are converted to YOLO format by calculating the normalized center coordinates, width, and height as follows:

$$x_{\text{center}} = \frac{x_{\min} + x_{\max}}{2 \times \text{image width}}, \quad y_{\text{center}} = \frac{y_{\min} + y_{\max}}{2 \times \text{image height}}$$

$$\text{width} = \frac{x_{\max} - x_{\min}}{\text{image width}}, \quad \text{height} = \frac{y_{\max} - y_{\min}}{\text{image height}}$$

Label File Format

Each line in the label file corresponds to one detected object and includes the class ID and normalized bounding box parameters in the format:

```
<class_id> <x_center> <y_center> <width> <height>
```

Handling No Detections

Even if no objects are detected, the script creates an empty label file to maintain consistency in the dataset structure.

This approach is commonly used in preparing datasets for YOLO training, especially when you want to generate labels automatically from an existing model before fine-tuning or retraining on your custom data.

```

1 from ultralytics import YOLO
2 import glob
3 import os
4
5 # Load the pre-trained YOLOv8 model
6 model = YOLO('yolov8n.pt')
7
8 # Define your custom dataset folder
9 image_dir = '/Sanjeev/VUNIT_CLASSES/VUNIT-AAI-SEM3/custom_dataset_yolo/train_1'
10 image_paths = glob.glob(os.path.join(image_dir, '*'))
11
12 # Run inference
13 results = model.predict(
14     image_paths,
15     save=False, # Don't save images with bounding boxes
16     verbose=False # Suppress verbose output
17 )
18
19 # Process each result
20 for i, result in enumerate(results):
21     image_path = image_paths[i] # Get the corresponding image path
22     label_file_name = os.path.splitext(os.path.basename(image_path))[0] + '.txt',

```

```

23     label_file_path = os.path.join(image_dir, label_file_name)
24
25     with open(label_file_path, 'w') as f:
26         has_objects = False
27         for box in result.bounding_boxes:
28             xyxy = box.xyxy.cpu().numpy()[0]
29             conf = box.conf.cpu().numpy()[0]
30             cls = int(box.cls.cpu().numpy()[0])
31             label = model.names[cls]
32
33             # Convert bounding box to YOLO format
34             x_center = (xyxy[0] + xyxy[2]) / (2 * result.orig_shape[1])
35             y_center = (xyxy[1] + xyxy[3]) / (2 * result.orig_shape[0])
36             width = (xyxy[2] - xyxy[0]) / result.orig_shape[1]
37             height = (xyxy[3] - xyxy[1]) / result.orig_shape[0]
38
39             # Write to label file
40             f.write(f"{cls} {x_center:.6f} {y_center:.6f} {width:.6f} {height:.6f}\n")
41             has_objects = True
42
43         if not has_objects:
44             print(f"No objects detected in {image_path}, created empty label file.")
45         else:
46             print(f"Label file created for {image_path}")

```

Listing 5: Write YOLO Label and create bounding box txt file

1.3.3 custom_data.yaml

This YAML file is a configuration file used by YOLO (You Only Look Once) object detection models (such as YOLOv5, YOLOv8, etc.) to specify the dataset details for training and validation. It tells the model where to find images and labels, how many classes to expect, and what the class names are.

```

1 train: /Sanjeev/VNIT_CLASSES/VNIT-AAI-SEM3/custom_dataset_yolo/train_1
2 val: /Sanjeev/VNIT_CLASSES/VNIT-AAI-SEM3/custom_dataset_yolo/test_1
3
4 nc: 20 # number of classes
5 names: ['bear', 'bird', 'bougainvillea', 'cat', 'cow', 'daisie', 'deer', 'dog', 'dolphin', 'elephant', 'ros

```

Listing 6: Write YOLO Label and create bounding box txt file

1.3.4 programming_assignment_q3_3.py

Purpose This script uses the Ultralytics YOLOv8 model to:

Run object detection inference on a set of images.

Save the inference results (images with bounding boxes).

Print detected object labels, bounding box coordinates, and confidence scores.

Evaluate the model on your custom dataset using the custom_data.yaml configuration and print key performance metrics (precision, recall, mAP).

```

1 from ultralytics import YOLO
2 import glob
3
4 # 1. Load the pre-trained YOLOv8 model
5 model = YOLO('yolov8n.pt')
6
7 # 2. Train the model on your custom dataset for 50 epochs
8 model.train(data='custom_data.yaml', epochs=50)
9

```

```

10 # 3. List of image paths for inference
11 image_paths = glob.glob('custom_dataset/*')
12
13 # 4. Run inference and save results to results/3_inference/
14 results = model.predict(
15     image_paths,
16     save=True,                                     # Save images with bounding boxes
17     project='results',                            # Top-level results folder
18     name='3_inference',                           # Subfolder for this run
19     exist_ok=True                                # Overwrite if the folder exists
20 )
21
22 # 5. Print detected boxes, labels, and confidence scores
23 for result in results:
24     for box in result.boxes:
25         xyxy = box.xyxy.cpu().numpy()[0]           # bounding box coordinates
26         conf = box.conf.cpu().numpy()[0]             # confidence score
27         cls = int(box.cls.cpu().numpy()[0])          # class index
28         label = model.names[cls]                   # class label
29         print(f"Label: {label}, Confidence: {conf:.2f}, Box: {xyxy}")
30
31 # 6. Run validation on your dataset YAML file
32 metrics = model.val(data='custom_data.yaml')
33
34 # 7. Print evaluation metrics safely
35 try:
36     print(f"2. Recall: {metrics.box.mr():.4f}")
37     print(f"2. mAP@0.5 (map50): {metrics.box.map50():.4f}")
38     print(f"4. mAP@0.5:0.95 (map): {metrics.box.map():.4f}")
39     print(f"1. Precision: {metrics.box.mp():.4f}")
40 except Exception as e:
41     print("Metrics could not be calculated. Check if your validation set has labels.")
42     print(f"Error: {e}")

```

Listing 7: Object detection inference

```

1 C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\subjects\com\anamika_singh\YOL01_Assignment>python programming_assignment.py
2 New https://pypi.org/project/ultralytics/8.3.141 available Update with 'pip install -U ultralytics'
3 Ultralytics 8.3.139 Python-3.11.9 torch-2.7.0+cpu CPU (Intel Core(TM) i7-8650U 1.90GHz)
4 engine\trainer: agnostic_nms=False, amp=True, augment=False, auto_augment=randaugment, batch=16, bgr=0.0, b
5 Overriding model.yaml nc=80 with nc=20
6
7          from    n      params   module
8  0          -1    1        464  ultralytics.nn.modules.conv.Conv
9  1          -1    1       4672  ultralytics.nn.modules.conv.Conv
10 2          -1    1       7360  ultralytics.nn.modules.block.C2f
11 3          -1    1      18560  ultralytics.nn.modules.conv.Conv
12 4          -1    2      49664  ultralytics.nn.modules.block.C2f
13 5          -1    1      73984  ultralytics.nn.modules.conv.Conv
14 6          -1    2     197632  ultralytics.nn.modules.block.C2f
15 7          -1    1     295424  ultralytics.nn.modules.conv.Conv
16 8          -1    1     460288  ultralytics.nn.modules.block.C2f
17 9          -1    1     164608  ultralytics.nn.modules.block.SPPF
18 10         -1    1        0  torch.nn.modules.upsampling.Upsample
19 11         [-1, 6]  1        0  ultralytics.nn.modules.conv.Concat
20 12         -1    1    148224  ultralytics.nn.modules.block.C2f
21 13         -1    1        0  torch.nn.modules.upsampling.Upsample
22 14         [-1, 4]  1        0  ultralytics.nn.modules.conv.Concat
23 15         -1    1     37248  ultralytics.nn.modules.block.C2f
24 16         -1    1     36992  ultralytics.nn.modules.conv.Conv
25 17         [-1, 12] 1        0  ultralytics.nn.modules.conv.Concat
26 18         -1    1    123648  ultralytics.nn.modules.block.C2f
27 19         -1    1    147712  ultralytics.nn.modules.conv.Conv
28 20         [-1, 9]  1        0  ultralytics.nn.modules.conv.Concat
29 21         -1    1    493056  ultralytics.nn.modules.block.C2f
          arguments
          [3, 16, 3, 2]
          [16, 32, 3, 2]
          [32, 32, 1, True]
          [32, 64, 3, 2]
          [64, 64, 2, True]
          [64, 128, 3, 2]
          [128, 128, 2, True]
          [128, 256, 3, 2]
          [256, 256, 1, True]
          [256, 256, 5]
          [None, 2, 'nearest']
          [1]
          [384, 128, 1]
          [None, 2, 'nearest']
          [1]
          [192, 64, 1]
          [64, 64, 3, 2]
          [1]
          [192, 128, 1]
          [128, 128, 3, 2]
          [1]
          [384, 256, 1]

```

```

30 22      [15, 18, 21] 1    755212  ultralytics.nn.modules.head.Detect          [20, [64, 128, 256]]
31 Model summary: 129 layers, 3,014,748 parameters, 3,014,732 gradients, 8.2 GFLOPs
32
33 Transferred 319/355 items from pretrained weights
34 Freezing layer 'model.22.dfl.conv.weight'
35 train: Fast image access (ping: 0.40.0 ms, read: 3.23.8 MB/s, size: 38.4 KB)
36 train: Scanning C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1... 98 images, 16 background
37 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Bear_17.jpg: ignoring corrupt image
38 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Bear_21_1.jpg: ignoring corrupt image
39 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Bear_3_2.jpg: ignoring corrupt image
40 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Bear_6.jpg: ignoring corrupt image
41 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Bird_4_3.jpg: ignoring corrupt image
42 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Bird_7_2.jpg: ignoring corrupt image
43 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Elephant_23_4.jpg: ignoring corrupt image
44 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Elephant_28_2.jpg: ignoring corrupt image
45 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Giraffe_21.jpg: ignoring corrupt image
46 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Giraffe_21_1.jpg: ignoring corrupt image
47 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Giraffe_28_4.jpg: ignoring corrupt image
48 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Giraffe_6_2.jpg: ignoring corrupt image
49 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Lion_6.jpg: ignoring corrupt image
50 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Panda_18.jpg: ignoring corrupt image
51 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Panda_20_1.jpg: ignoring corrupt image
52 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Panda_23.jpg: ignoring corrupt image
53 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Panda_25_2.jpg: ignoring corrupt image
54 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Panda_3_2.jpg: ignoring corrupt image
55 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Tiger_27_4.jpg: ignoring corrupt image
56 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Tiger_8.jpg: ignoring corrupt image
57 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Zebra_28_1.jpg: ignoring corrupt image
58 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Zebra_30_3.jpg: ignoring corrupt image
59 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\Zebra_3_1.jpg: ignoring corrupt image
60 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\bear_1.jpg: ignoring corrupt image
61 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\bougainvillea_00004.jpg: ignoring corrupt image
62 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\cat_189.jpg: ignoring corrupt image
63 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\daisies_00022.jpg: ignoring corrupt image
64 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\dog_109.jpg: ignoring corrupt image
65 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\dog_260.jpg: ignoring corrupt image
66 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\gardenias_00026.jpg: ignoring corrupt image
67 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\gardenias_00051.jpg: ignoring corrupt image
68 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\hydrangeas_00013.jpg: ignoring corrupt image
69 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\hydrangeas_00026.jpg: ignoring corrupt image
70 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\img_p1_11.jpg: ignoring corrupt image
71 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\img_p1_111.jpg: ignoring corrupt image
72 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\img_p1_20appapp.jpg: ignoring corrupt image
73 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\img_p1_72.jpg: ignoring corrupt image
74 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\img_p1_92.jpg: ignoring corrupt image
75 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\img_p2_100.jpg: ignoring corrupt image
76 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\img_p2_69.jpg: ignoring corrupt image
77 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\img_p3_46.jpg: ignoring corrupt image
78 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\img_p3_58.jpg: ignoring corrupt image
79 train: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1\img_p3_70appapp.jpg: ignoring corrupt image
80 train: New cache created: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\train_1.cache
81 C:\Users\urssa\AppData\Local\Programs\Python\Python311\Lib\site-packages\torch\utils\data\dataloader.py:665
82     warnings.warn(warn_msg)
83 val: Fast image access (ping: 0.40.0 ms, read: 7.36.3 MB/s, size: 71.8 KB)
84 val: Scanning C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1.cache... 28 images, 9 background
85 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\Bear_10.jpg: ignoring corrupt image
86 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\Bird_15.jpg: ignoring corrupt image
87 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\Zebra_20.jpg: ignoring corrupt image
88 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\dog_442.jpg: ignoring corrupt image
89 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\hydrangeas_00070.jpg: ignoring corrupt image
90 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\img_p1_25.jpg: ignoring corrupt image
91 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\img_p3_124.jpg: ignoring corrupt image
92 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\img_p3_50.jpg: ignoring corrupt image
93 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\lilies_00069.jpg: ignoring corrupt image
94 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\peonies_00006.jpg: ignoring corrupt image

```

```

95 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\peonies_00077.jpg: ignoring corrupt i
96 C:\Users\urssa\AppData\Local\Programs\Python\Python311\Lib\site-packages\torch\utils\data\dataloader.py:665
97 warnings.warn(warn_msg)
98 Plotting labels to runs\detect\train\labels.jpg...
99 optimizer: 'optimizer=auto' found, ignoring 'lr0=0.01' and 'momentum=0.937' and determining best 'optimizer'
100 optimizer: AdamW(lr=0.000417, momentum=0.9) with parameter groups 57 weight(decay=0.0), 64 weight(decay=0.0
101 Image sizes 640 train, 640 val
102 Using 0 dataloader workers
103 Logging results to runs\detect\train
104 Starting training for 50 epochs...
105
106   Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
107   1/50      OG        0.7892     4.191      1.341       11          640: 100%|
108           Class    Images    Instances    Box(P)      R          mAP50    mAP50-95): 100%|
109           all         17          9          0          0          0          0
110
111   Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
112   2/50      OG        0.7977     4.271      1.364       16          640: 100%|
113           Class    Images    Instances    Box(P)      R          mAP50    mAP50-95): 100%|
114           all         17          9          0          0          0          0
115
116   Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
117   3/50      OG        0.7394     4.112      1.286       20          640: 100%|
118           Class    Images    Instances    Box(P)      R          mAP50    mAP50-95): 100%|
119           all         17          9          0          0          0          0
120
121   Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
122   4/50      OG        0.6739     3.966      1.267       19          640: 100%|
123           Class    Images    Instances    Box(P)      R          mAP50    mAP50-95): 100%|
124           all         17          9          0          0          0          0
125
126   Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
127   5/50      OG        0.6346     3.782      1.219       15          640: 100%|
128           Class    Images    Instances    Box(P)      R          mAP50    mAP50-95): 100%|
129           all         17          9          0          0          0          0
130
131   Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
132   6/50      OG        0.6584     3.85       1.238       16          640: 100%|
133           Class    Images    Instances    Box(P)      R          mAP50    mAP50-95): 100%|
134           all         17          9          0          0          0          0
135
136   Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
137   7/50      OG        0.6517     3.715      1.281       16          640: 100%|
138           Class    Images    Instances    Box(P)      R          mAP50    mAP50-95): 100%|
139           all         17          9          0.1        0.0667      0.1        0.0801
140
141   Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
142   8/50      OG        0.6487     3.555      1.215       19          640: 100%|
143           Class    Images    Instances    Box(P)      R          mAP50    mAP50-95): 100%|
144           all         17          9          0.111      0.233      0.209        0.148
145
146   Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
147   9/50      OG        0.6167     3.436      1.181       22          640: 100%|
148           Class    Images    Instances    Box(P)      R          mAP50    mAP50-95): 100%|
149           all         17          9          0.0461      0.233      0.202        0.15
150
151   Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
152   10/50     OG        0.6602     3.328      1.129       16          640: 100%|
153           Class   Images    Instances    Box(P)      R          mAP50    mAP50-95): 100%|
154           all         17          9          0.0565      0.5          0.236        0.177
155
156   Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
157   11/50     OG        0.6446     3.091      1.19        18          640: 100%|
158           Class   Images    Instances    Box(P)      R          mAP50    mAP50-95): 100%|
159           all         17          9          0.05        0.5          0.186        0.147

```

160								
161	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
162	12/50	OG	0.5791	2.995	1.151	16	640: 100%	
163		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
164		all	17	9	0.124	0.8	0.309	0.229
165								
166	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
167	13/50	OG	0.682	2.936	1.266	17	640: 100%	
168		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
169		all	17	9	0.0938	1	0.392	0.29
170								
171	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
172	14/50	OG	0.7101	2.981	1.256	8	640: 100%	
173		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
174		all	17	9	0.0833	1	0.45	0.35
175								
176	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
177	15/50	OG	0.6377	2.794	1.261	17	640: 100%	
178		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
179		all	17	9	0.0518	1	0.54	0.417
180								
181	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
182	16/50	OG	0.6553	2.485	1.209	20	640: 100%	
183		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
184		all	17	9	0.509	0.433	0.522	0.409
185								
186	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
187	17/50	OG	0.6873	2.591	1.261	24	640: 100%	
188		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
189		all	17	9	0.638	0.297	0.472	0.366
190								
191	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
192	18/50	OG	0.6339	2.472	1.178	24	640: 100%	
193		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
194		all	17	9	0.639	0.277	0.465	0.353
195								
196	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
197	19/50	OG	0.6937	2.544	1.276	19	640: 100%	
198		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
199		all	17	9	0.861	0.267	0.427	0.308
200								
201	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
202	20/50	OG	0.6271	2.389	1.198	17	640: 100%	
203		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
204		all	17	9	0.878	0.267	0.428	0.322
205								
206	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
207	21/50	OG	0.6518	2.245	1.186	12	640: 100%	
208		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
209		all	17	9	0.973	0.323	0.487	0.351
210								
211	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
212	22/50	OG	0.6789	2.044	1.264	16	640: 100%	
213		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
214		all	17	9	0.961	0.333	0.49	0.359
215								
216	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
217	23/50	OG	0.655	2.201	1.206	14	640: 100%	
218		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
219		all	17	9	0.934	0.333	0.47	0.343
220								
221	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
222	24/50	OG	0.6205	2.23	1.185	16	640: 100%	
223		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%
224		all	17	9	0.915	0.333	0.464	0.328

225	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
226	25/50	OG	0.5582	2.058	1.185	17	640: 100%
227		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
228		all	17	9	0.656	0.333	0.483 0.348
229							
230	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
231	26/50	OG	0.6237	1.98	1.16	14	640: 100%
232		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
233		all	17	9	0.649	0.38	0.53 0.367
234							
235	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
236	27/50	OG	0.5437	1.979	1.124	19	640: 100%
237		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
238		all	17	9	0.791	0.433	0.607 0.402
239							
240	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
241	28/50	OG	0.5789	1.829	1.127	17	640: 100%
242		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
243		all	17	9	0.832	0.433	0.604 0.426
244							
245	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
246	29/50	OG	0.6148	1.904	1.174	21	640: 100%
247		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
248		all	17	9	0.769	0.433	0.598 0.421
249							
250	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
251	30/50	OG	0.5341	1.741	1.112	22	640: 100%
252		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
253		all	17	9	0.658	0.433	0.693 0.458
254							
255	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
256	31/50	OG	0.7069	2.187	1.247	12	640: 100%
257		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
258		all	17	9	0.576	0.433	0.599 0.413
259							
260	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
261	32/50	OG	0.571	1.905	1.149	15	640: 100%
262		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
263		all	17	9	0.632	0.433	0.596 0.409
264							
265	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
266	33/50	OG	0.5401	1.713	1.118	17	640: 100%
267		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
268		all	17	9	0.733	0.433	0.696 0.475
269							
270	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
271	34/50	OG	0.5495	1.796	1.091	19	640: 100%
272		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
273		all	17	9	0.847	0.432	0.695 0.467
274							
275	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
276	35/50	OG	0.5194	1.648	1.098	17	640: 100%
277		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
278		all	17	9	0.867	0.421	0.509 0.369
279							
280	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
281	36/50	OG	0.5542	1.734	1.115	25	640: 100%
282		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
283		all	17	9	0.875	0.411	0.49 0.362
284							
285	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
286	37/50	OG	0.5231	1.67	1.084	15	640: 100%
287		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100%
288		all	17	9	0.803	0.367	0.449 0.35
289							

```

290
291     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
292     38/50    OG          0.5477    1.696      1.103      15          640: 100%|
293           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
294           all        17         9          0.803      0.367      0.446       0.35
295
296     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
297     39/50    OG          0.5714    1.886      1.161      12          640: 100%|
298           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
299           all        17         9          0.803      0.367      0.444       0.356
300
301     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
302     40/50    OG          0.5699    1.708      1.193      11          640: 100%|
303           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
304           all        17         9          0.802      0.367      0.502       0.383
305 Closing dataloader mosaic
306 C:\Users\urssa\AppData\Local\Programs\Python\Python311\Lib\site-packages\torch\utils\data\dataloader.py:665
307     warnings.warn(warn_msg)
308
309     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
310     41/50    OG          0.4338    2.588      1.125      3          640: 100%|
311           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
312           all        17         9          0.808      0.367      0.593       0.427
313
314     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
315     42/50    OG          0.4441    2.404      1.01       9          640: 100%|
316           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
317           all        17         9          0.79       0.367      0.461       0.362
318
319     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
320     43/50    OG          0.473     2.291      1.136      11         640: 100%|
321           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
322           all        17         9          0.768      0.367      0.476       0.37
323
324     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
325     44/50    OG          0.3864    2.104      1.001      9          640: 100%|
326           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
327           all        17         9          0.747      0.367      0.505       0.385
328
329     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
330     45/50    OG          0.4293    2.288      0.9662     8          640: 100%|
331           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
332           all        17         9          0.734      0.366      0.502       0.38
333
334     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
335     46/50    OG          0.4818    2.316      1.158       7          640: 100%|
336           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
337           all        17         9          0.726      0.359      0.504       0.386
338
339     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
340     47/50    OG          0.4392    2.486      1.087       6          640: 100%|
341           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
342           all        17         9          0.721      0.353      0.498       0.381
343
344     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
345     48/50    OG          0.4401    2.375      1.169       6          640: 100%|
346           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
347           all        17         9          0.735      0.36       0.504       0.388
348
349     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size
350     49/50    OG          0.4808    2.273      1.126       6          640: 100%|
351           Class    Images    Instances    Box(P)      R          mAP50   mAP50-95): 100%|
352           all        17         9          0.749      0.362      0.511       0.391
353
354     Epoch    GPU_mem    box_loss    cls_loss    dfl_loss    Instances    Size

```

```

355      50/50      OG      0.4778      2.336      1.149       7      640: 100%|
356          Class     Images   Instances     Box(P)        R      mAP50  mAP50-95): 100%|
357          all       17         9       0.758      0.363      0.478      0.373
358
359 50 epochs completed in 0.594 hours.
360 Optimizer stripped from runs\detect\train\weights\last.pt, 6.3MB
361 Optimizer stripped from runs\detect\train\weights\best.pt, 6.3MB
362
363 Validating runs\detect\train\weights\best.pt...
364 Ultralytics 8.3.139 Python-3.11.9 torch-2.7.0+cpu CPU (Intel Core(TM) i7-8650U 1.90GHz)
365 Model summary (fused): 72 layers, 3,009,548 parameters, 0 gradients, 8.1 GFLOPs
366          Class     Images   Instances     Box(P)        R      mAP50  mAP50-95): 100%|
367          all       17         9       0.736      0.433      0.696      0.475
368          bear      1         1         1         0       0.995      0.497
369          apple     2         2         1         0       0.19       0.145
370          tomato    2         3       0.726      0.667      0.72       0.579
371          kangaroo  2         2       0.375      0.5       0.578      0.357
372          lion      1         1       0.58       1       0.995      0.796
373 Speed: 4.8ms preprocess, 189.9ms inference, 0.0ms loss, 2.8ms postprocess per image
374 Results saved to runs\detect\train
375
376 0: 640x640 (no detections), 143.2ms
377 1: 640x640 (no detections), 143.2ms
378 2: 640x640 (no detections), 143.2ms
379 3: 640x640 (no detections), 143.2ms
380 4: 640x640 (no detections), 143.2ms
381 5: 640x640 (no detections), 143.2ms
382 6: 640x640 2 tomatoes, 143.2ms
383 7: 640x640 (no detections), 143.2ms
384 Speed: 6.6ms preprocess, 143.2ms inference, 2.8ms postprocess per image at shape (1, 3, 640, 640)
385 Results saved to results\3_inference
386 Label: tomato, Confidence: 0.87, Box: [      5.0869      21.313     1482.3      1192]
387 Label: tomato, Confidence: 0.34, Box: [     46.278      25.329      1064     1179.2]
388 Ultralytics 8.3.139 Python-3.11.9 torch-2.7.0+cpu CPU (Intel Core(TM) i7-8650U 1.90GHz)
389 val: Fast image access (ping: 0.30.0 ms, read: 63.365.1 MB/s, size: 68.3 KB)
390 val: Scanning C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1.cache... 28 images, 9 backgr
391 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\Bear_10.jpg: ignoring corrupt image/1
392 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\Bird_15.jpg: ignoring corrupt image/1
393 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\Zebra_20.jpg: ignoring corrupt image/1
394 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\dog_442.jpg: ignoring corrupt image/1
395 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\hydrangeas_00070.jpg: ignoring corrupt image/1
396 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\img_p1_25.jpeg: ignoring corrupt image/1
397 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\img_p3_124.jpeg: ignoring corrupt image/1
398 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\img_p3_50.jpeg: ignoring corrupt image/1
399 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\lilies_00069.jpg: ignoring corrupt image/1
400 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\peonies_00006.jpg: ignoring corrupt image/1
401 val: C:\Sanjeev\VNIT_CLASSES\VNIT-AAI-SEM3\custom_dataset_yolo\test_1\peonies_00077.jpg: ignoring corrupt image/1
402 C:\Users\urssa\AppData\Local\Programs\Python\Python311\Lib\site-packages\torch\utils\data\dataloader.py:668
403     warnings.warn(warn_msg)
404          Class     Images   Instances     Box(P)        R      mAP50  mAP50-95): 100%|
405          all       17         9       0.736      0.433      0.696      0.475
406          bear      1         1         1         0       0.995      0.497
407          apple     2         2         1         0       0.19       0.145
408          tomato    2         3       0.726      0.667      0.72       0.579
409          kangaroo  2         2       0.375      0.5       0.578      0.357
410          lion      1         1       0.58       1       0.995      0.796
411 Speed: 8.3ms preprocess, 336.5ms inference, 0.0ms loss, 2.6ms postprocess per image
412 Results saved to runs\detect\train2
413 Metrics could not be calculated. Check if your validation set has labels.
414 Error: 'numpy.float64' object is not callable

```

Listing 8: Object detection inference execution output

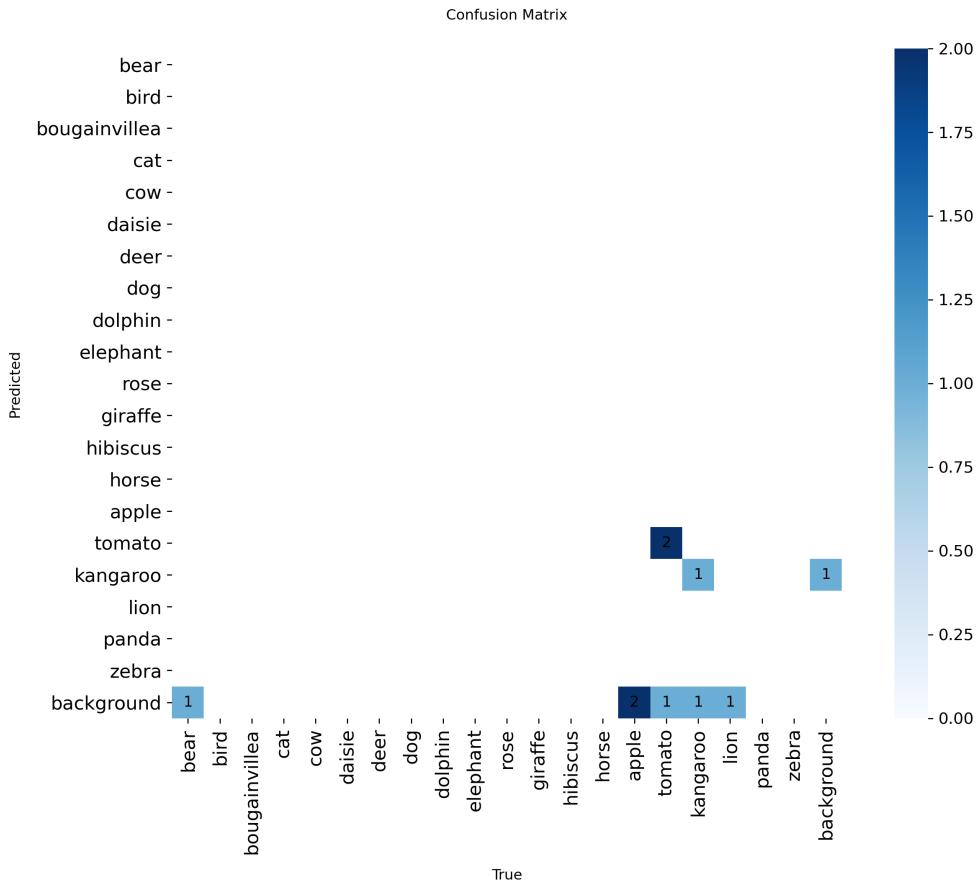


Figure 9: Confusion Matrix

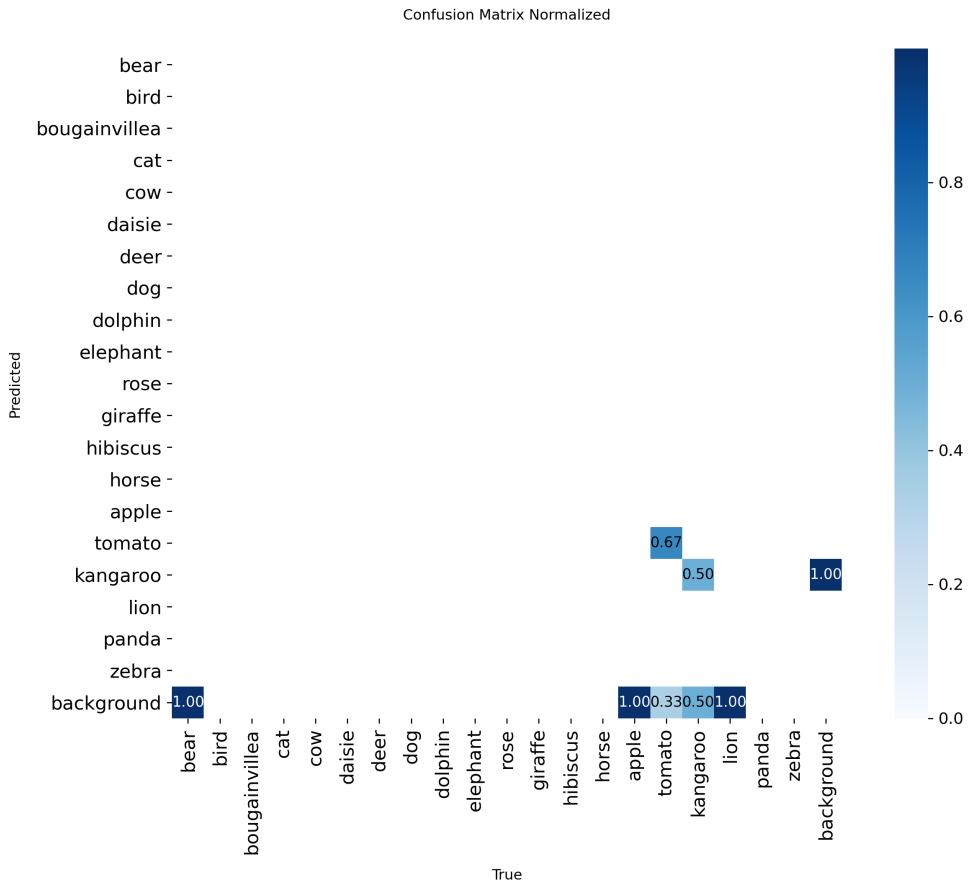


Figure 10: Normalized Confusion Matrix

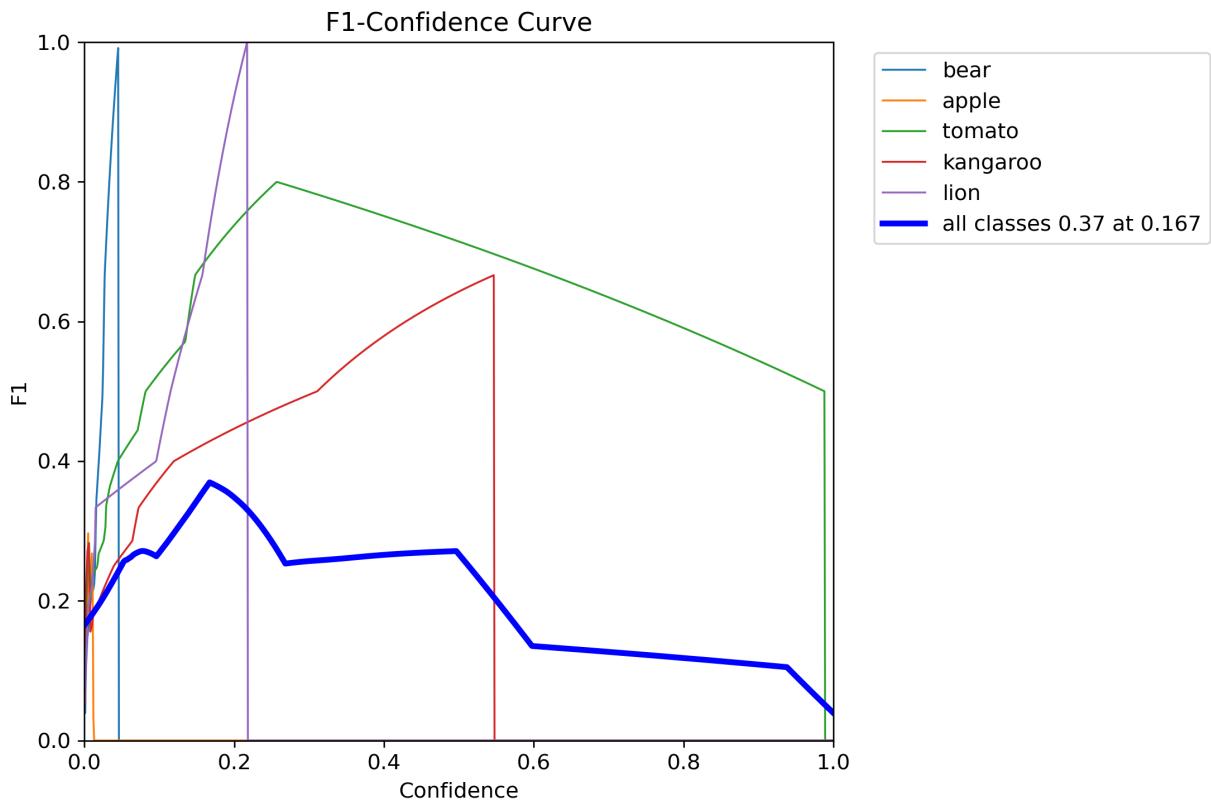


Figure 11: F1 Score Curve

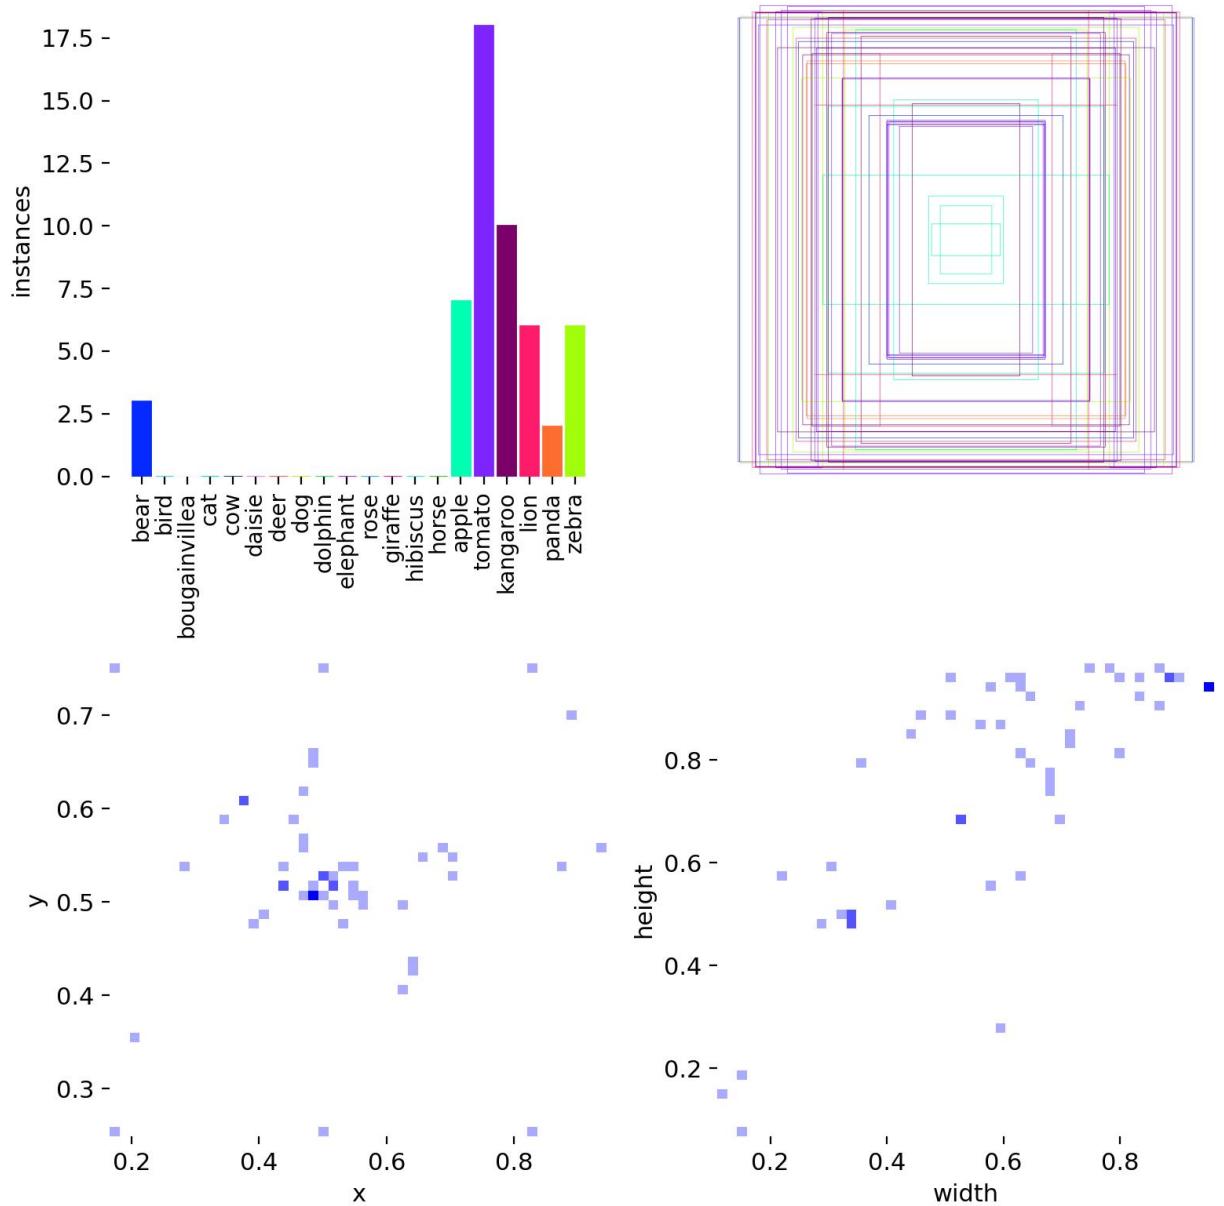


Figure 12: Labels Image

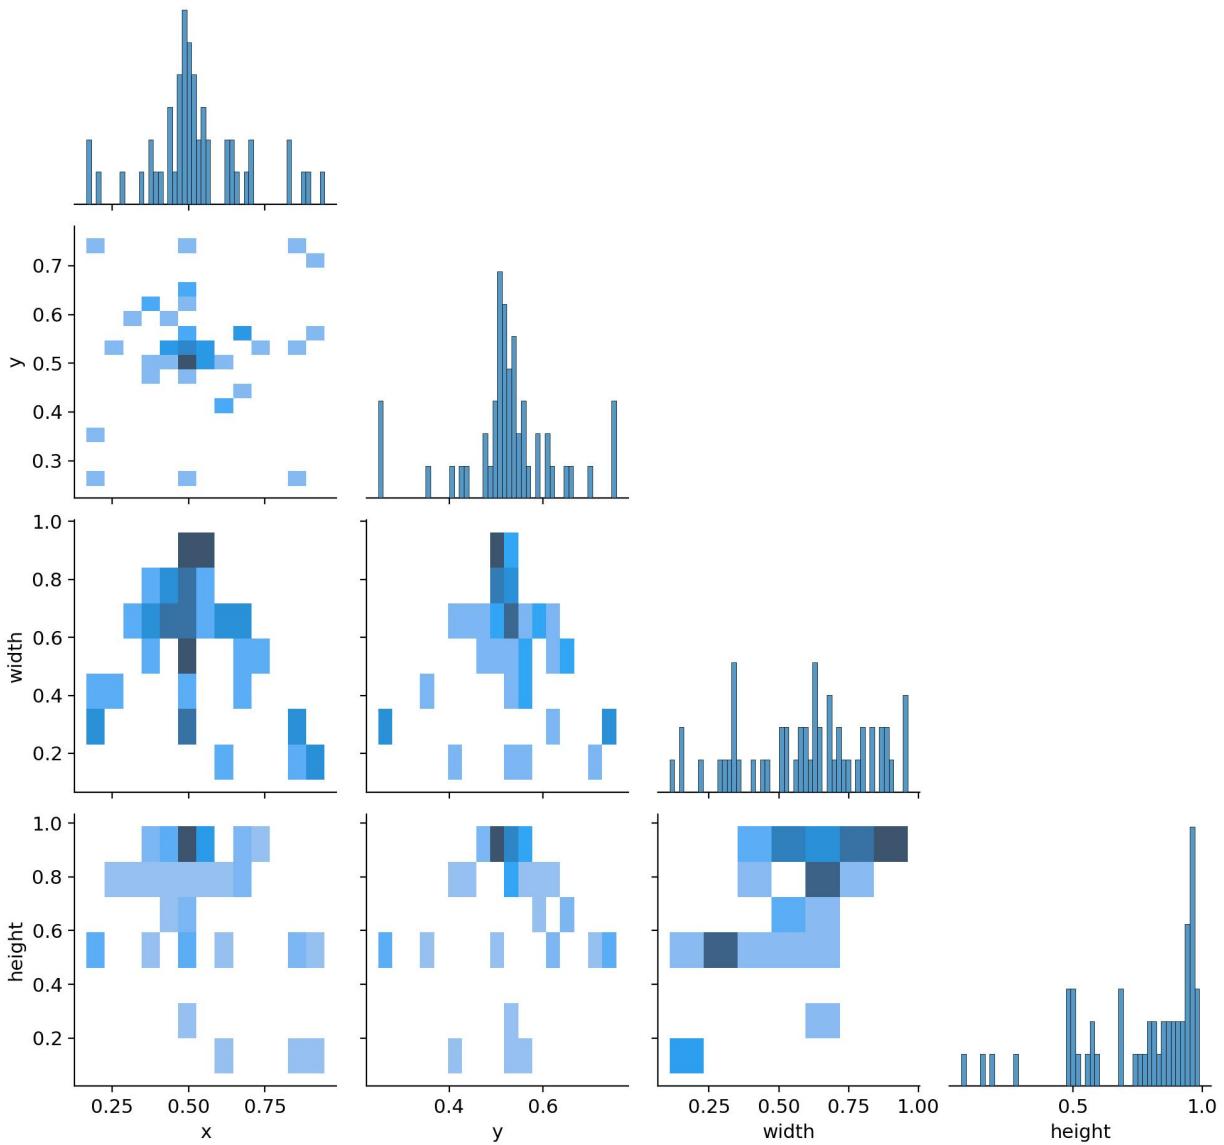


Figure 13: Labels Correlogram

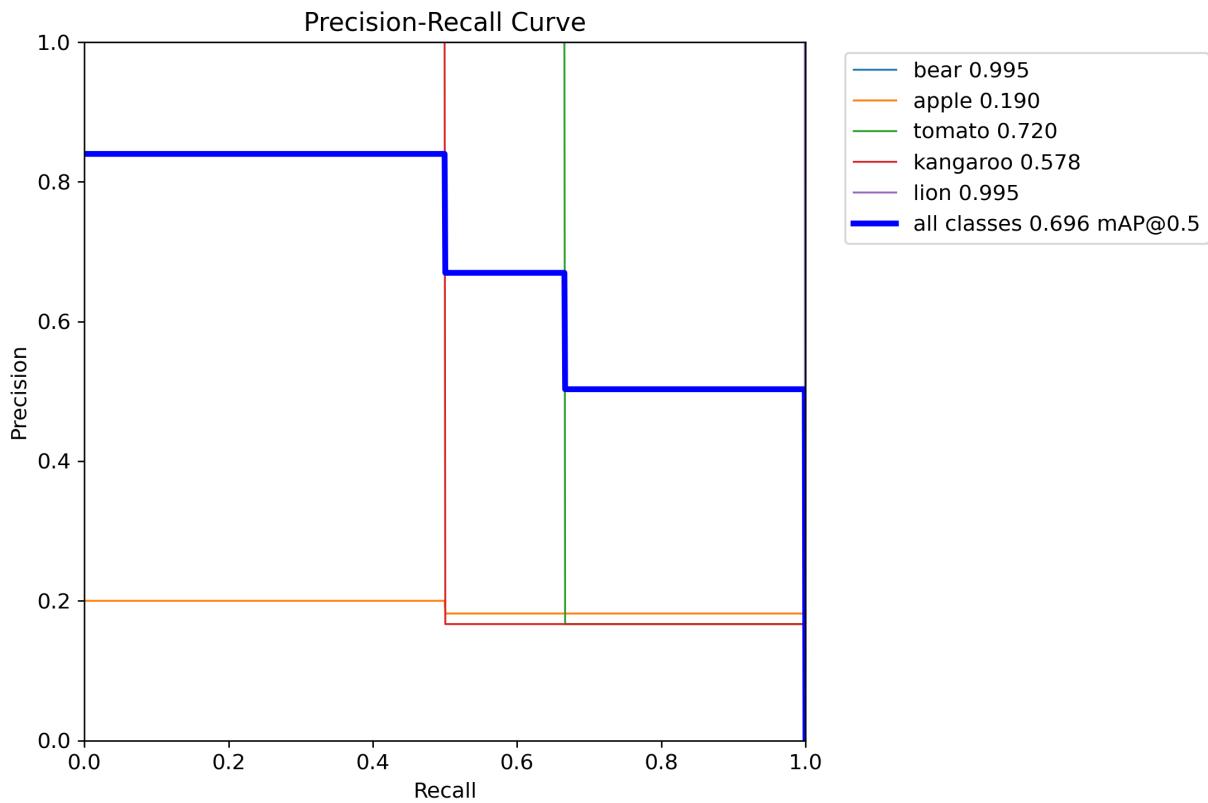


Figure 14: Precision-Recall Curve

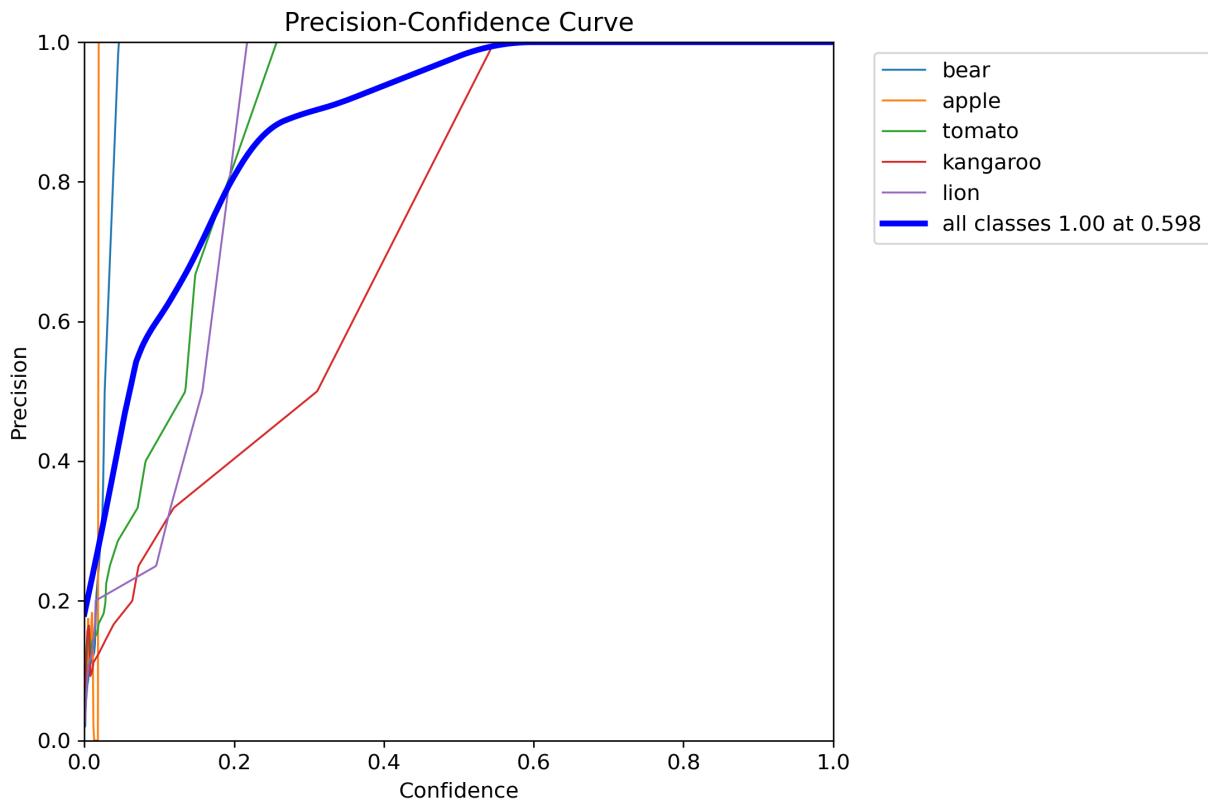


Figure 15: Precision Curve

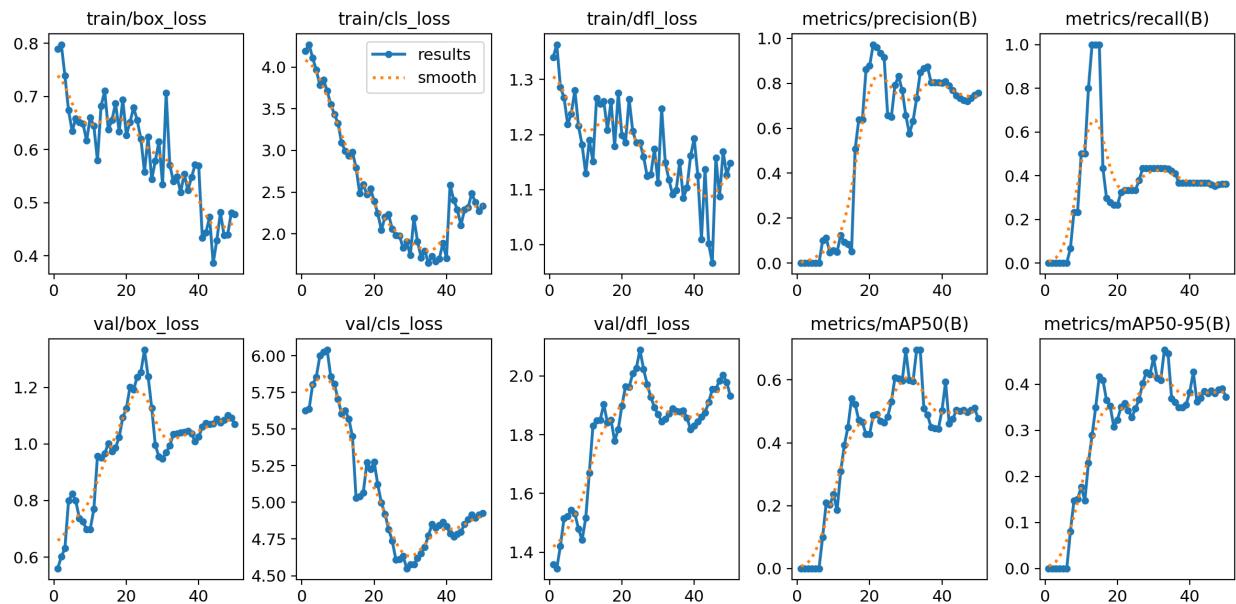


Figure 16: Results Image

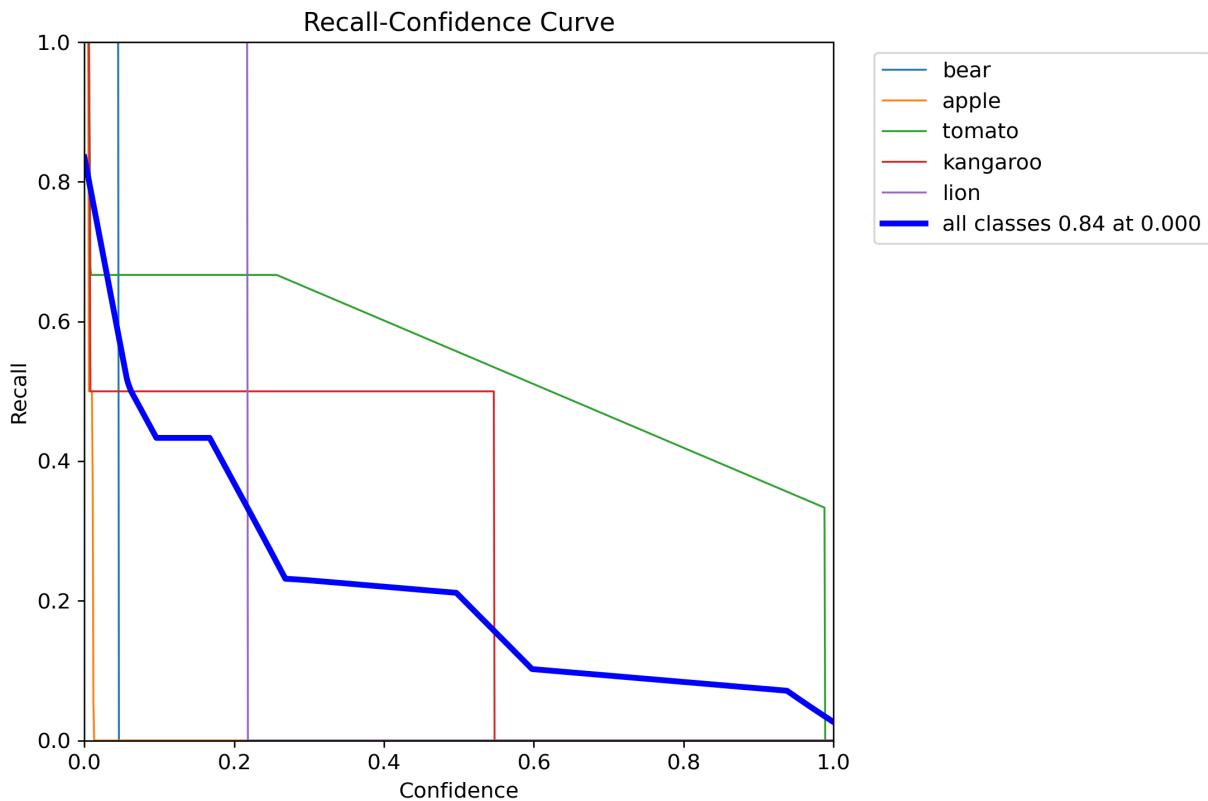


Figure 17: Recall Curve

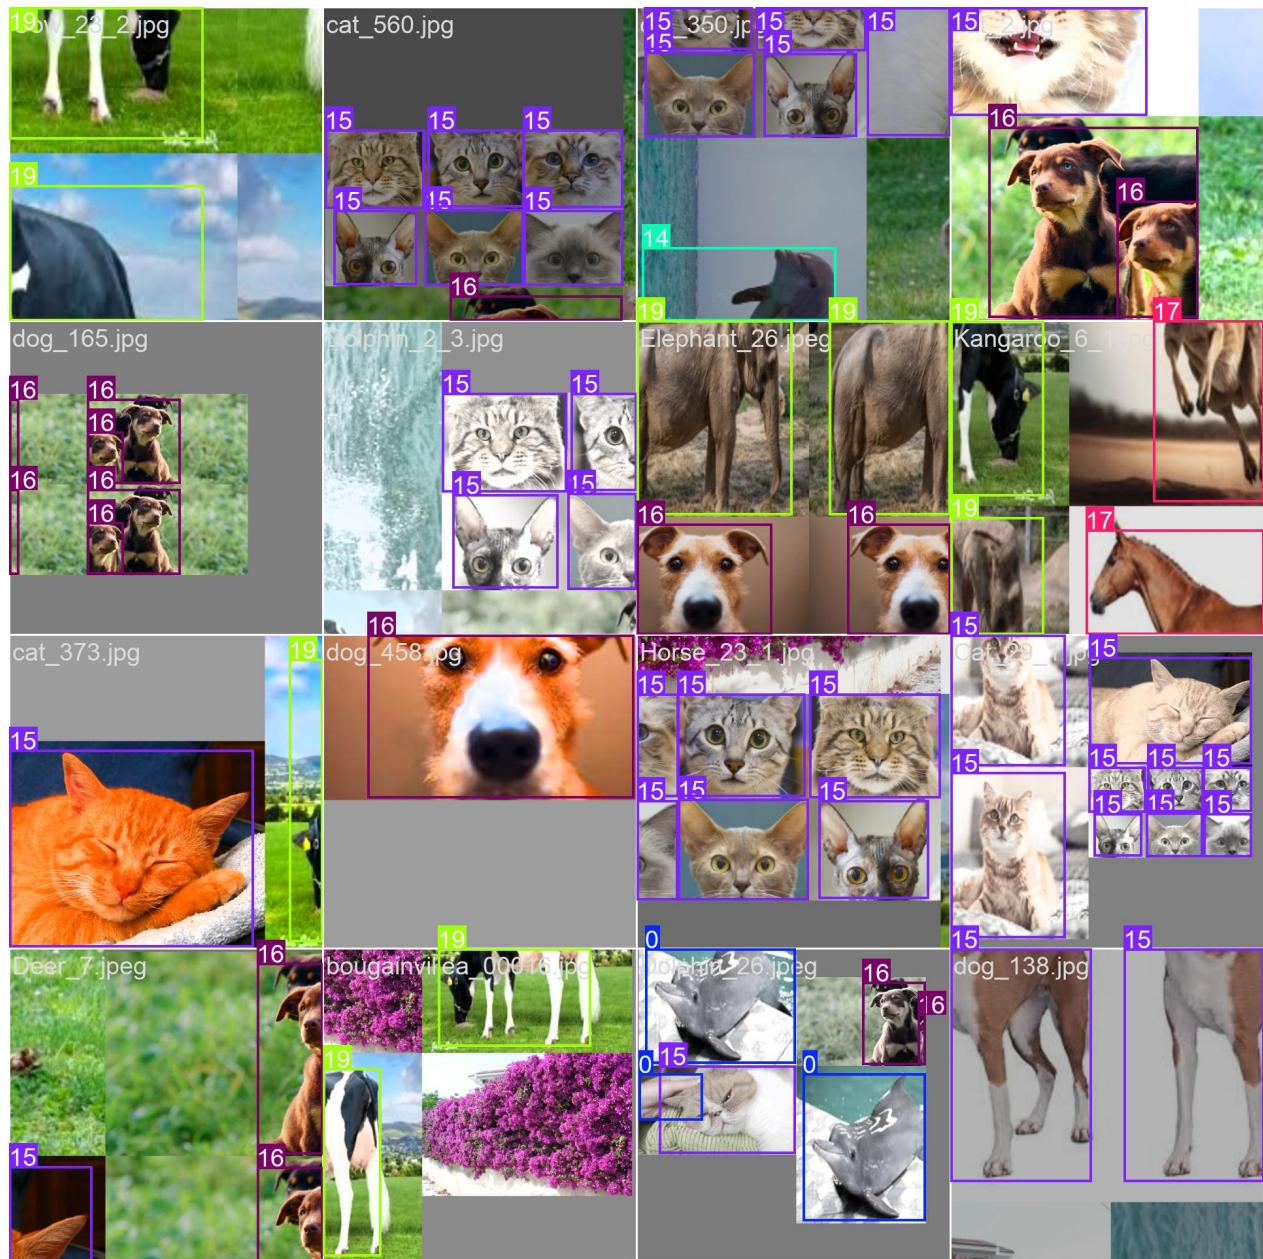


Figure 18: Training Batch 0

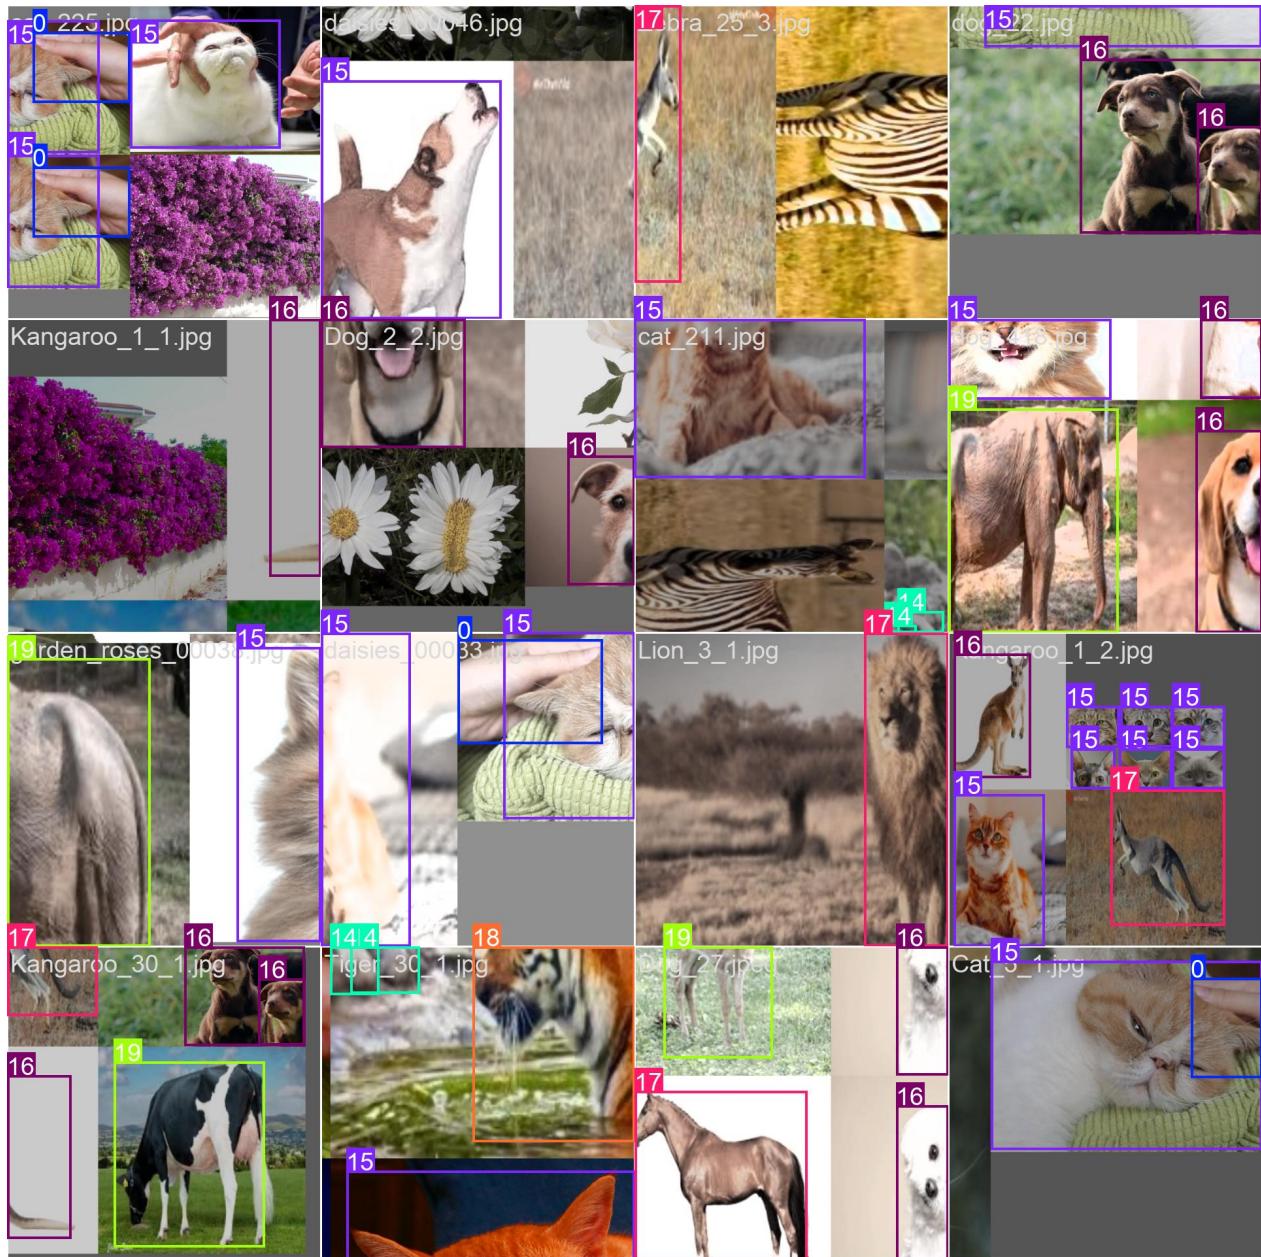


Figure 19: Training Batch 1

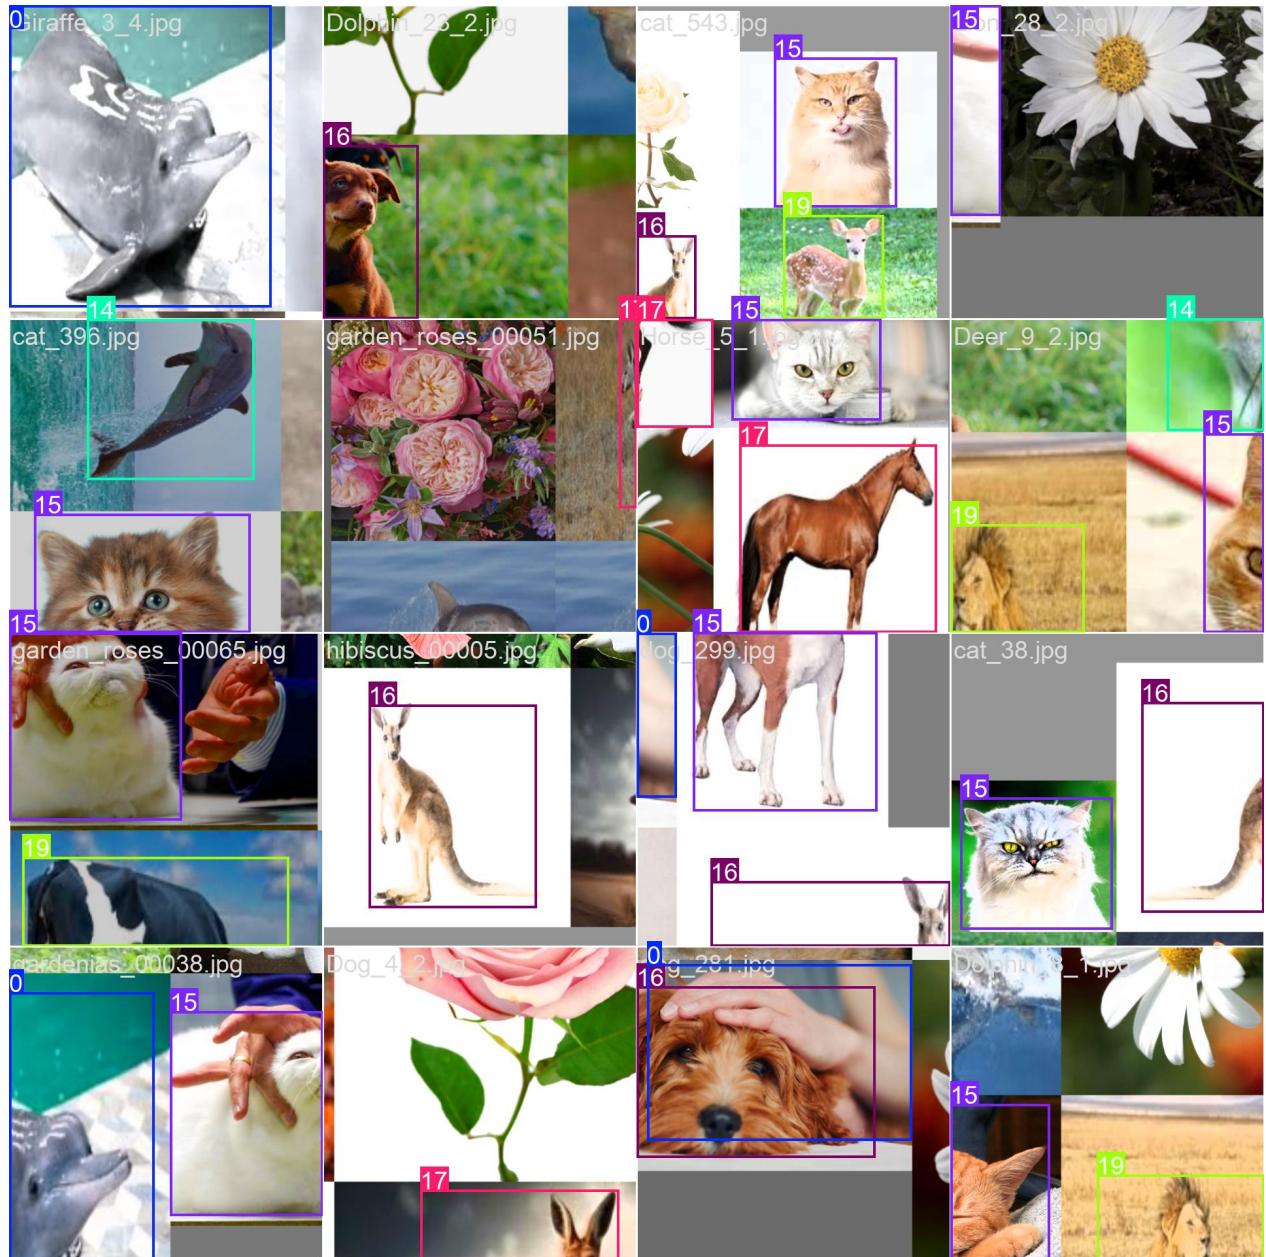


Figure 20: Training Batch 2

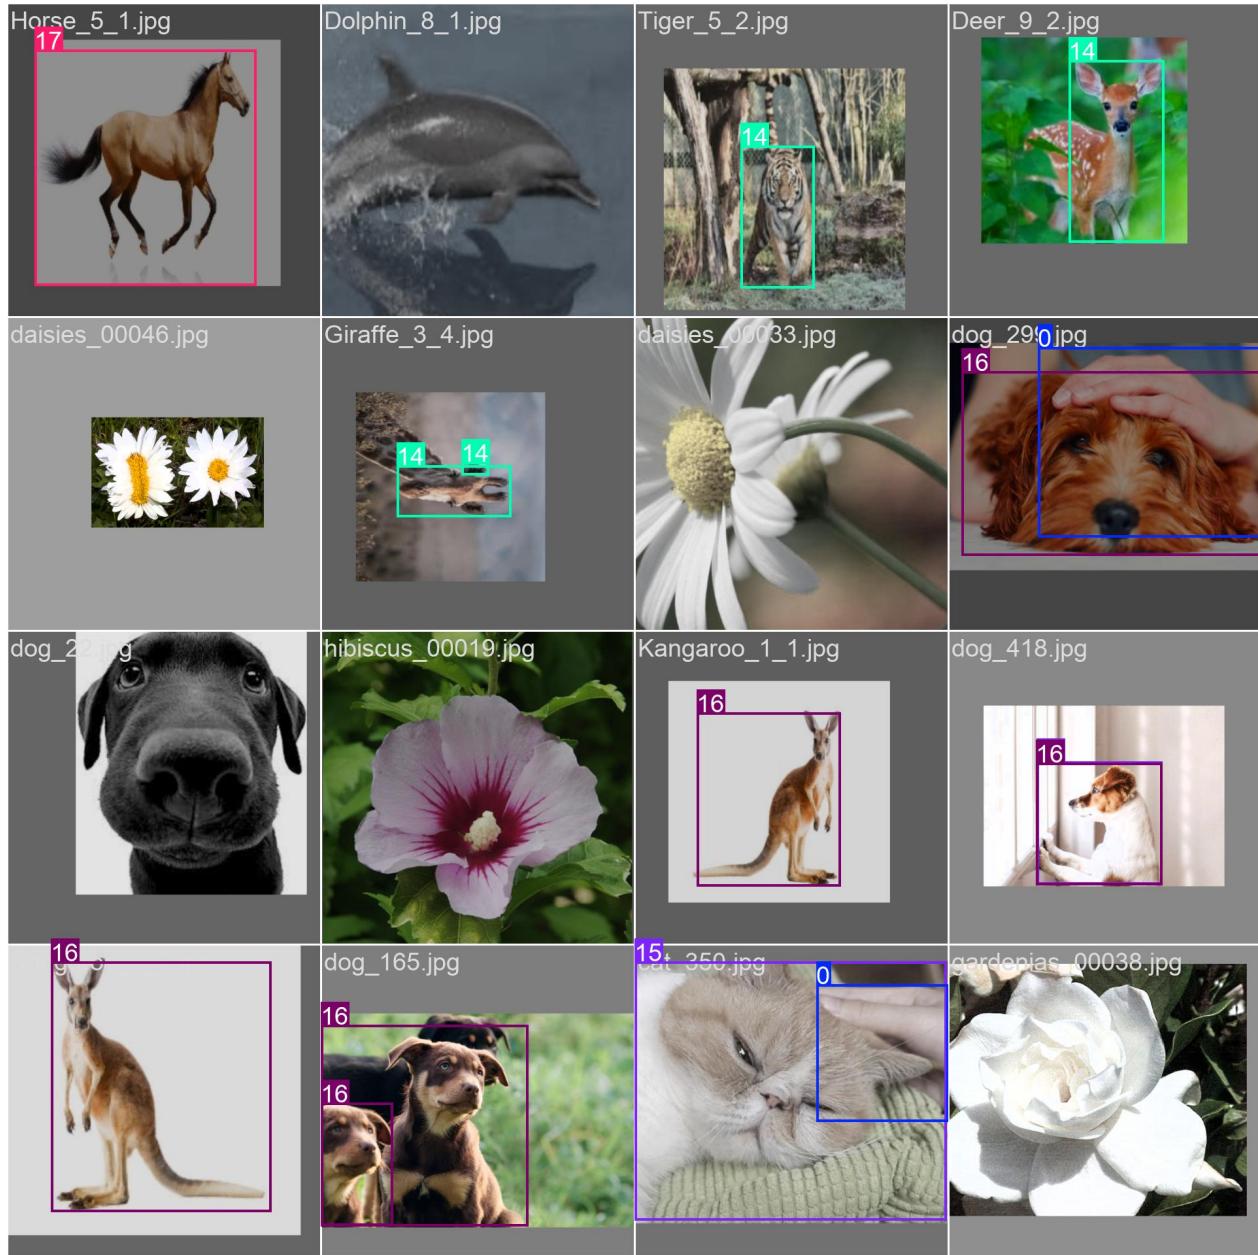


Figure 21: Training Batch 160

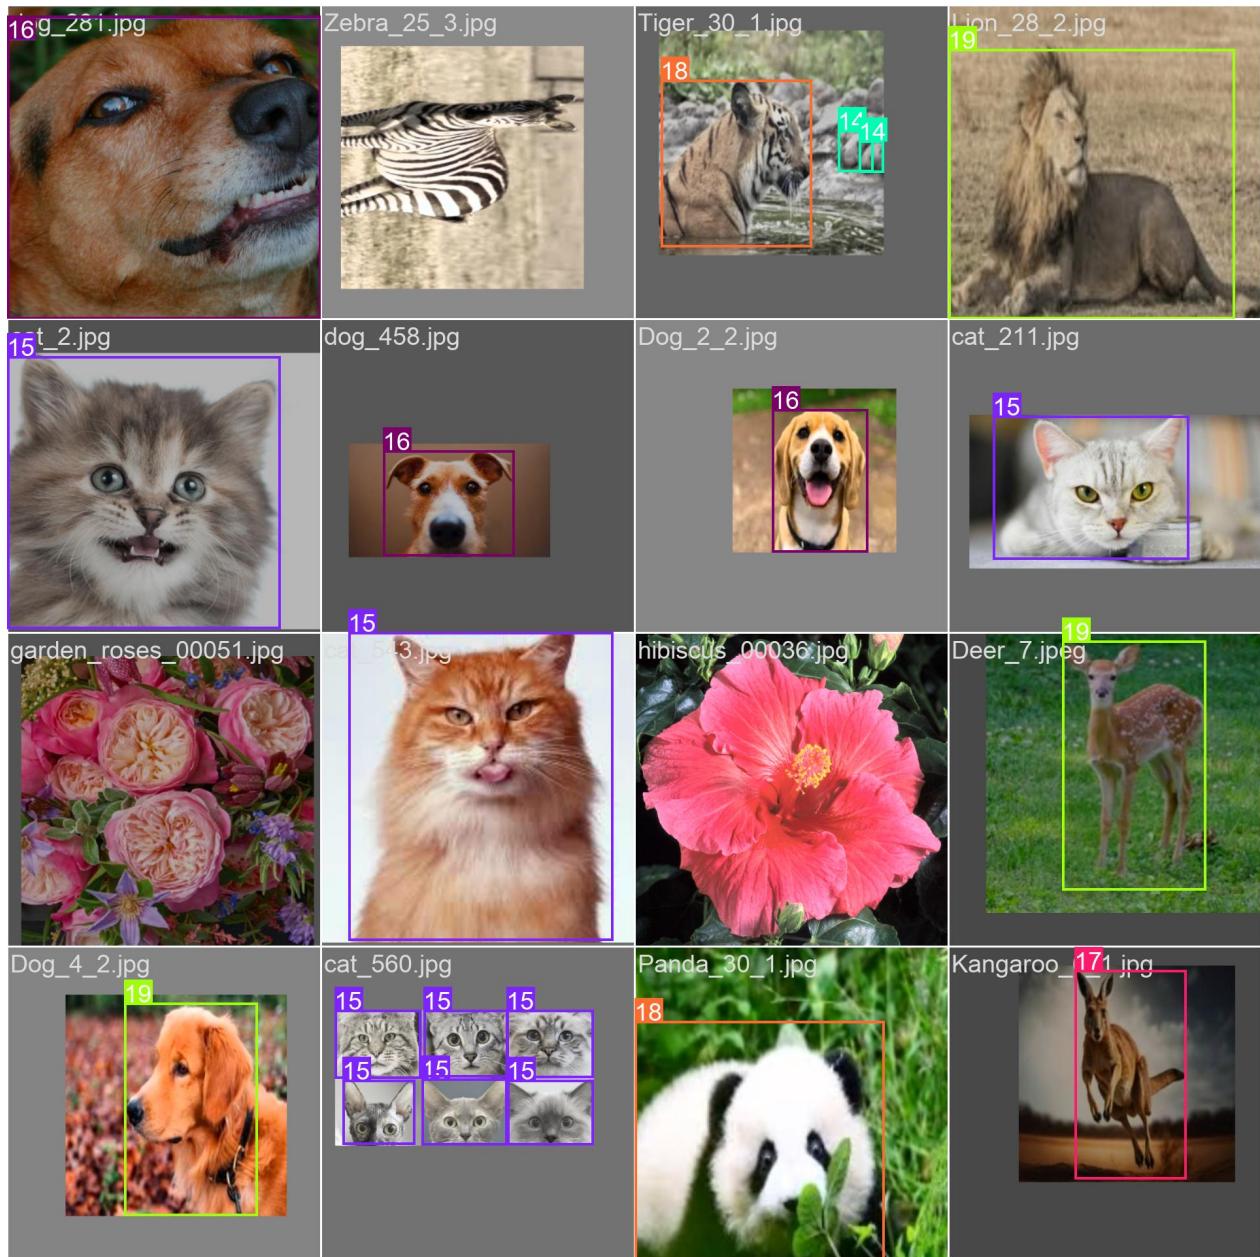


Figure 22: Training Batch 161

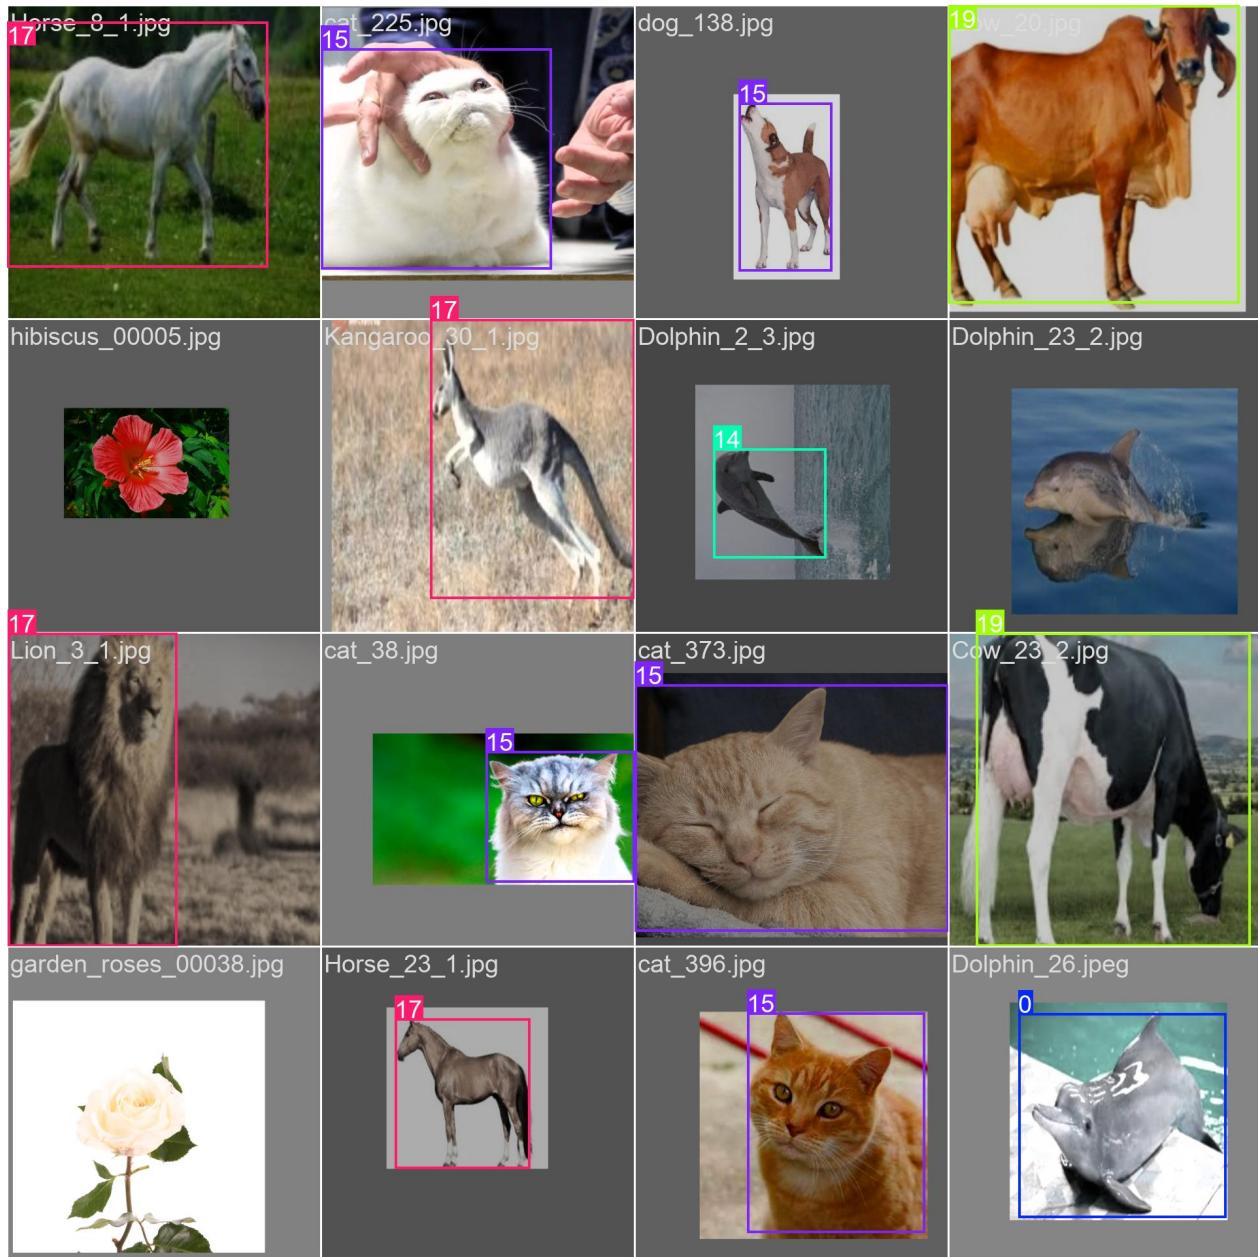


Figure 23: Training Batch 162

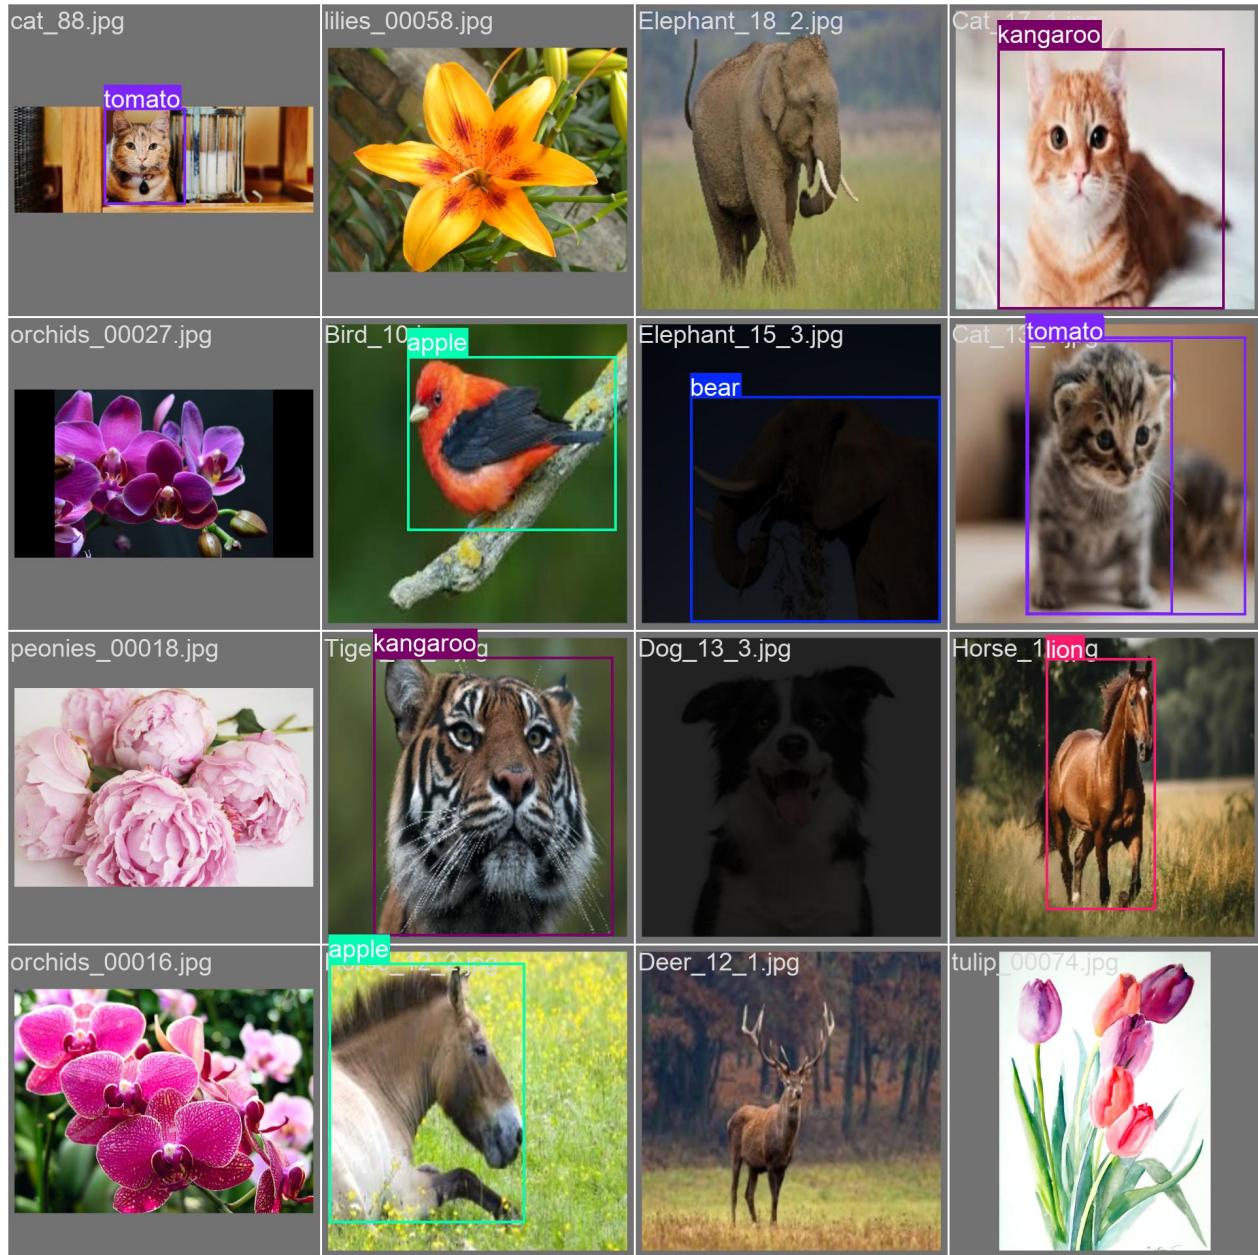


Figure 24: Validation Batch 0 Labels

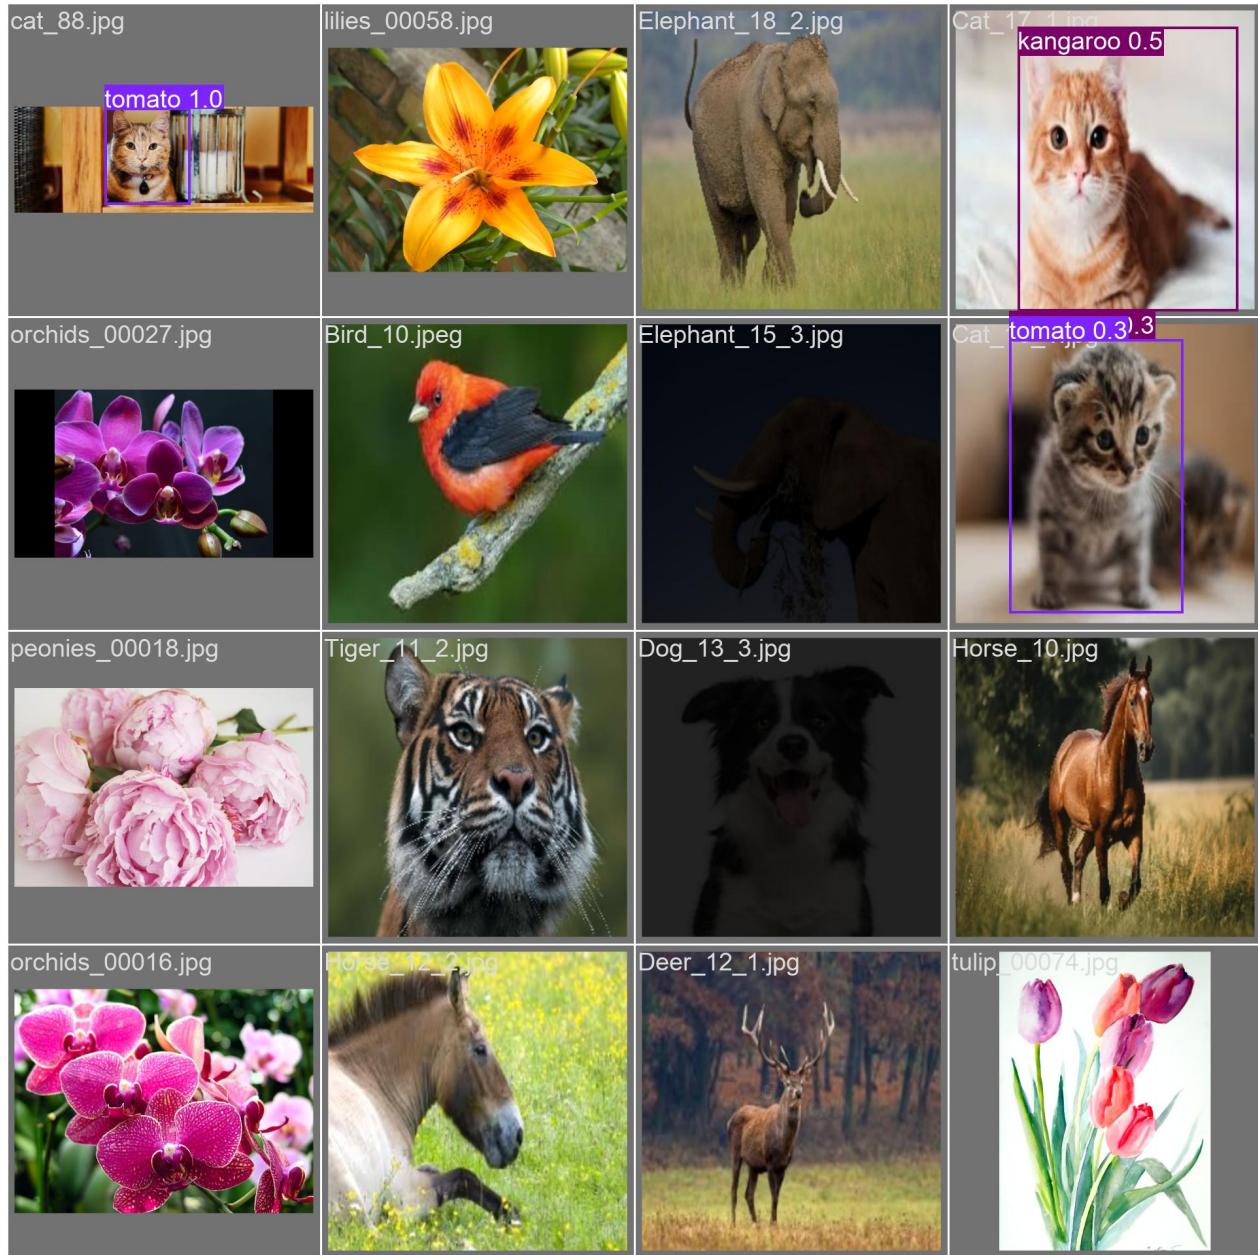


Figure 25: Validation Batch 0 Predictions

2 Theory Assignment

2.1 Question 1

How does YOLOv1 process input images? Describe the architecture of YOLOv1.

Answer

How YOLOv1 Processes Input Images and Its Architecture

YOLOv1 is a fast and clever way for a computer to find objects in pictures. Instead of looking at parts of the image one by one, YOLOv1 looks at the whole image just once and tries to find all the objects at the same time.

How It Works

First, YOLOv1 takes any input image and resizes it to a standard size so it can process it easily. Then, it divides this image into a grid made up of small squares. Each square is responsible for checking if there is an object inside it.

For each square, the model guesses a few boxes where objects might be, and also how confident it is about those guesses. It also predicts what kind of object is inside each box, like a dog, a car, or a person.

After making all these guesses for every square, YOLOv1 filters out overlapping boxes and keeps the best ones. This way, it quickly finds and labels objects in the image in just one go.

The Architecture

YOLOv1 uses a deep neural network made up of many layers that help it understand the image:

- It has a series of layers called convolutional layers that look for important features like edges and shapes.
- These layers shrink the image step-by-step while keeping the important details.
- After these, there are a couple of fully connected layers that take all the information from the previous layers and make the final predictions about where objects are and what they are.
- The network uses special functions that help it learn complex patterns from images.

Because the network looks at the entire image at once and predicts all objects in one step, YOLOv1 is very fast – fast enough to work in real-time applications like video surveillance or self-driving cars.

In Simple Terms

- YOLOv1 splits the image into a grid.
- Each grid cell guesses if there is an object inside it.
- It predicts boxes around objects and what those objects are.
- It does all this in a single pass, making it very fast.

This approach makes YOLOv1 different from older methods that needed to look at the image many times or in pieces. It's like glancing at a photo once and instantly knowing what's in it and where.

Detailed Architecture of YOLOv1

YOLOv1 is a single neural network designed to detect objects in images quickly and accurately by looking at the whole image at once. Its architecture is inspired by a well-known image classification network called GoogLeNet, but it has been adapted for the task of object detection.

Key Components of the Network

- **Input:** The network takes an image resized to 448 by 448 pixels as input, ensuring consistent processing regardless of the original image size.
- **Convolutional Layers:** YOLOv1 has **24 convolutional layers** that extract important features from the image. These layers scan the image for patterns such as edges, textures, and shapes that help identify objects.
- **Max-Pooling Layers:** Interspersed between convolutional layers are **max-pooling layers**, which reduce the spatial size of the feature maps. This helps the network focus on the most important information while reducing computational load.
- **1x1 and 3x3 Convolutions:** The network cleverly alternates between small 1x1 convolutions, which reduce the number of channels (making the network more efficient), and 3x3 convolutions, which capture spatial information.
- **Activation Functions:** Most layers use a **Leaky ReLU** activation function, which helps the network learn complex patterns by allowing small gradients when neurons are not active.
- **Fully Connected Layers:** After the convolutional layers, there are **2 fully connected layers** that take all the extracted features and make the final predictions. These layers output a fixed-size tensor containing bounding box coordinates, confidence scores, and class probabilities for each grid cell.
- **Output Tensor:** The output is structured as a grid (usually 7x7). Each grid cell predicts:
 - Multiple bounding boxes (usually 2 per cell), each with coordinates and a confidence score indicating how likely it is that the box contains an object.
 - Class probabilities indicating what type of object might be inside the bounding boxes.
- **Global Reasoning:** Because the fully connected layers see the entire image's features at once, the network can reason about the global context – understanding how objects relate to each other across the whole image.

Training Details and Techniques

- The network was initially pretrained on a large image classification dataset (ImageNet) to learn general visual features.
- It was then fine-tuned on object detection datasets like PASCAL VOC, where it learned to predict bounding boxes and classes.
- Techniques like **dropout** and **data augmentation** were used to prevent overfitting, helping the model generalize better to new images.
- The loss function combines errors from bounding box coordinates, confidence scores, and classification, training the network end-to-end as a single regression problem.

Why This Architecture Matters - The combination of convolutional layers and fully connected layers allows YOLOv1 to detect multiple objects in an image simultaneously and quickly. - The use of 1x1 convolutions reduces the number of parameters and speeds up computation without sacrificing accuracy. - Processing the entire image in one pass (instead of scanning parts individually) makes YOLOv1 much faster than previous object detection methods. - The network's design balances accuracy and speed, enabling real-time object detection on standard hardware.

Limitations to Keep in Mind - YOLOv1 can struggle with detecting very small objects because the grid cells are relatively large. - It can detect only a limited number of objects per image due to the fixed grid and bounding box predictions. - Nearby or overlapping objects can be hard to separate because each grid cell predicts only a few bounding boxes.

In summary, YOLOv1's architecture is a smart blend of deep convolutional layers and fully connected layers that work together to quickly and accurately detect objects across the entire image in one go. This design was groundbreaking at the time and laid the foundation for many improved versions of YOLO that followed.

2.2 Question 2

How has the architecture evolved in YOLO different versions over the last decade?

Answer

How YOLO Architecture Has Evolved Over Time

Over the past ten years, YOLO's design has changed a lot. Each new version brought improvements that help it find objects in images faster, more accurately, and in a wider variety of situations. Here's a simple overview of how YOLO's design has improved:

YOLOv1: The Beginning YOLOv1 introduced the idea of detecting objects by looking at the whole image once. It divided the image into a grid and predicted bounding boxes and object types for each grid cell all at once. It was fast and straightforward but had trouble with small objects and precise localization.

YOLOv2: Introducing Anchor Boxes YOLOv2 added anchor boxes, which helped it better detect objects of different shapes and sizes. It used a new backbone network called Darknet-19, which improved feature extraction. It also included batch normalization for faster, more stable training and trained on images of varying sizes to be more flexible.

YOLOv3: Multi-Scale Detection YOLOv3 upgraded to a deeper backbone called Darknet-53 and started making predictions at multiple scales, which helped it detect small objects better. It kept its speed while becoming much more accurate.

YOLOv4: Better Feature Fusion YOLOv4 introduced new techniques like Cross-Stage Partial connections (CSP) and Spatial Pyramid Pooling (SPP) to combine features more effectively. It also used a Path Aggregation Network (PAN) to improve information flow. These changes balanced speed and accuracy well.

YOLOv5: Easier to Use YOLOv5 switched to the PyTorch framework, making it easier to train and customize. It used a CSPDarknet53 backbone and offered different model sizes to balance speed and accuracy for different needs.

YOLOv6 and YOLOv7: Focus on Efficiency YOLOv6 introduced a RepVGG-style backbone for faster inference and better accuracy. YOLOv7 improved layer connections further, achieving top performance for real-time detection.

YOLOv8 and Beyond The latest versions continue refining the architecture and training methods to make the models faster, more accurate, and more adaptable to various tasks.

Version	Main Design Changes	Improvements
YOLOv1	Single CNN with grid-based predictions	Fast and simple detection
YOLOv2	Anchor boxes, Darknet-19 backbone, batch normalization	Better accuracy and flexibility
YOLOv3	Darknet-53 backbone, multi-scale predictions	Improved small object detection and accuracy
YOLOv4	CSP connections, SPP, PAN	Balanced speed and accuracy
YOLOv5	PyTorch implementation, CSPDarknet53, multiple model sizes	Easier training and deployment
YOLOv6	RepVGG backbone, improved blocks	Faster inference and higher accuracy
YOLOv7	Enhanced layer aggregation (E-ELAN)	State-of-the-art real-time performance
YOLOv8+	Lightweight, refined architecture	Better adaptability and deployment

In short, YOLOs design has grown from a simple, one-step detector into a series of highly optimized models. Each new version builds on the last to make object detection faster, more accurate, and easier to use in real-world applications.

3 End of Document